AWS re:Invent

ANT201

Big Data Analytics Architectural Patterns and Best Practices

Ben Snively
Solutions Architect
Amazon Web Services

AWS re:Invent

aws

What to expect from the session

- Big data challenges
- Architectural principles
- How to simplify big data processing
- What technologies should you use?
  - Why?
  - How?
- Reference architecture
- Design patterns
- Customer examples
- Demonstrations

Types of big data analytics

Batch/ interactive

Stream processing

Machine learning

Delivery of big data services

Virtualized

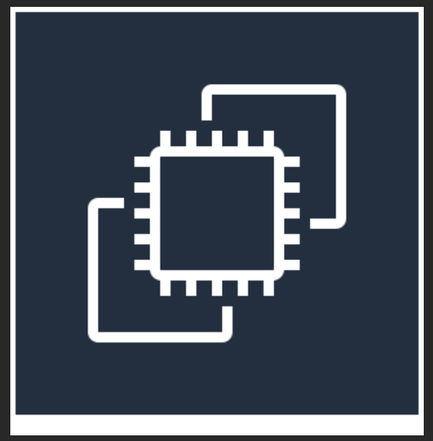

Managed services

Serverless/ clusterless/ containerized

Plethora of tools

AWS re:Invent

aws

Big data challenges

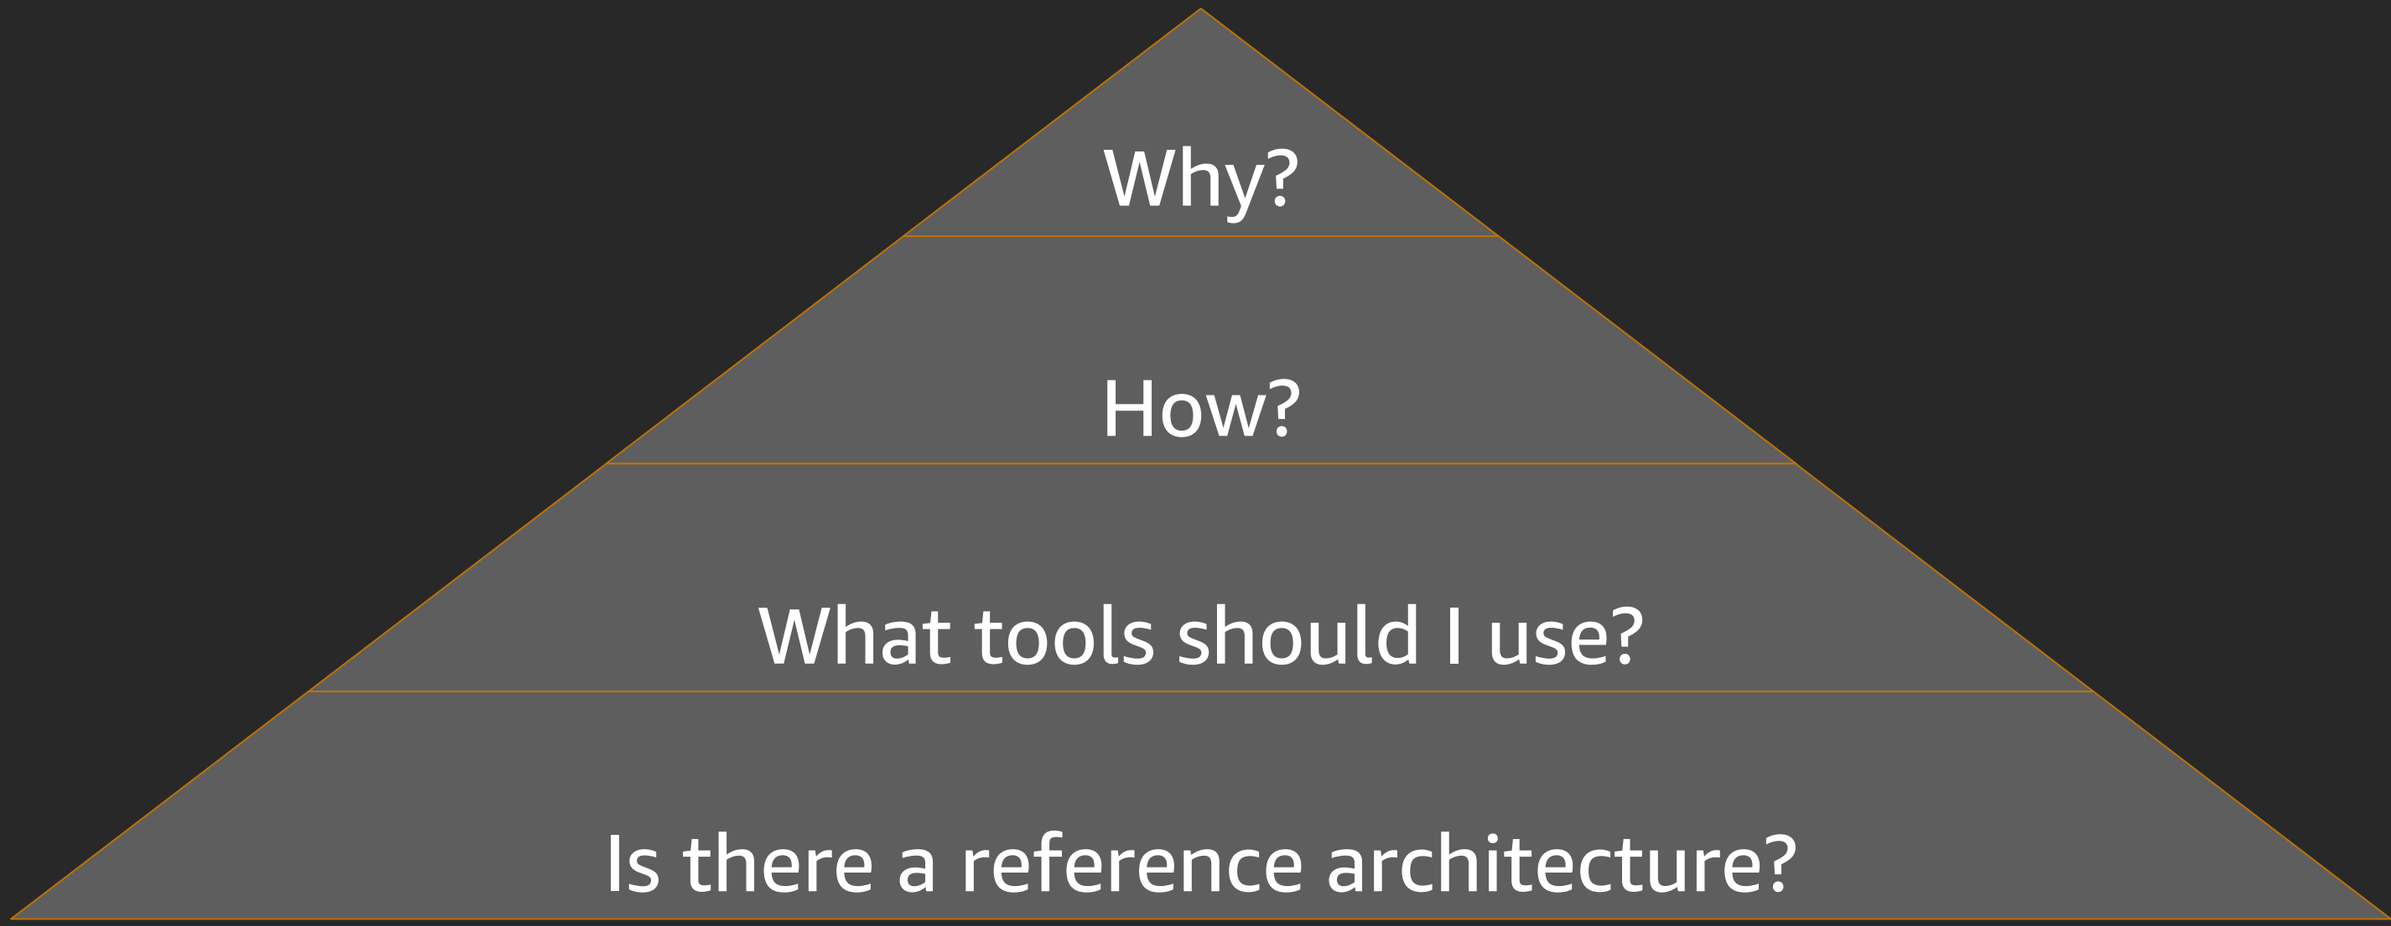

Architectural principles

Build decoupled systems

- Data → Store → Process → Store → Analyze → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Leverage managed and serverless services

- Scalable/elastic, available, reliable, secure, no/low admin

Use event-journal design patterns

- Immutable datasets (data lake), materialized views

Be cost-conscious

- Big data ≠ big cost

Machine learning (ML) enable your applications

Simplify big data processing

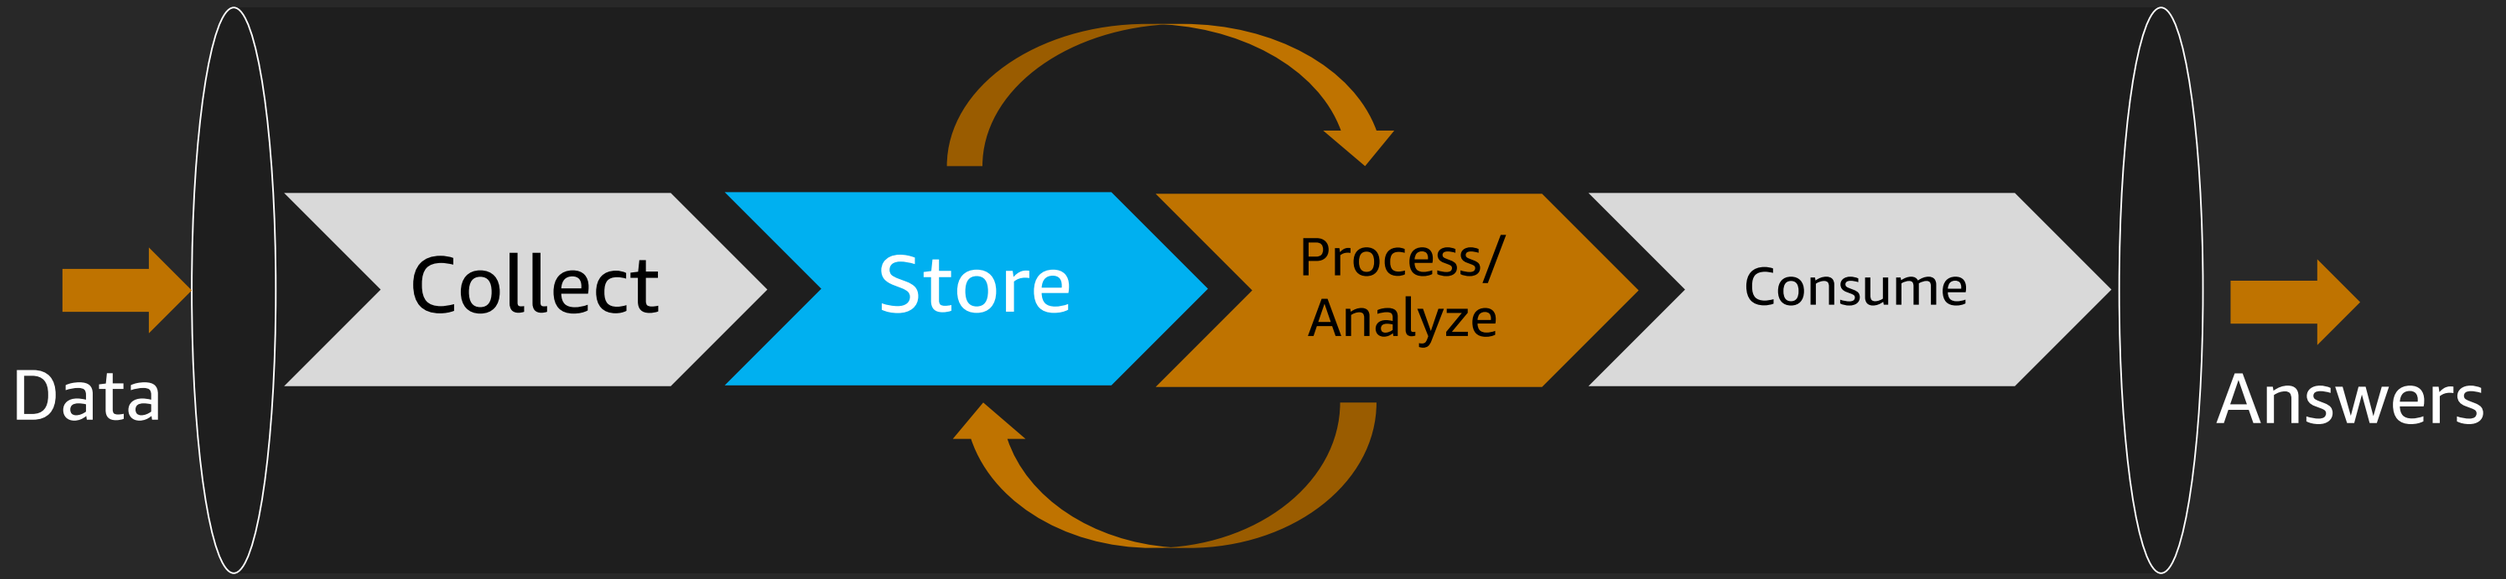

Time to answer (latency)
Throughput
Cost

What is the temperature of your data/analytic?

	Hot	Warm	Cold
Volume	MB–GB	GB–TB	PB–EB
Item size	B–KB	KB–MB	KB–TB
Latency	µs, ms	ms, sec	min, hrs
Durability	Low–high	High	Very high
Request rate	Very high	High	Low
Cost/GB	$$-$	$-¢¢	¢

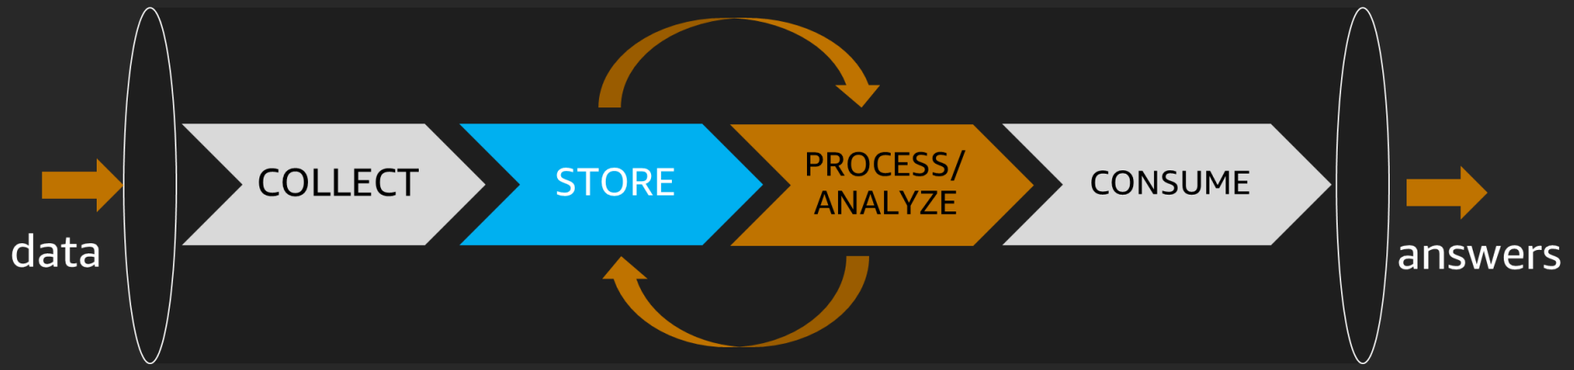

Collect

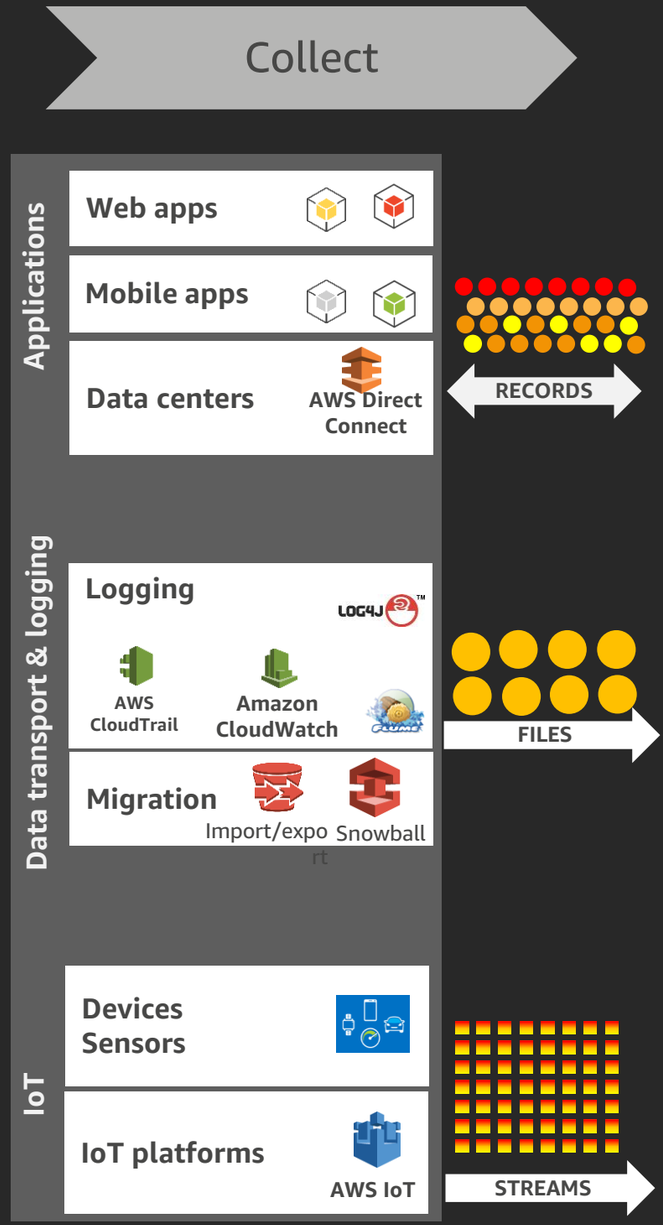

Type of data

Data structures

Database records

Transactions

Media files

Log files

Files/objects

Data streams

Events

aws

Type of data

Collect

Store

Applications

Web apps

Mobile apps

Data centers

AWS Direct Connect

RECORDS

Data structures

Database records

In-memory

NoSQL

SQL

Data transport & logging

Logging

LOG4J

AWS CloudTrail

Amazon CloudWatch

Migration

Import/export

Snowball

FILES

Media files

Log files

File/object store

IoT

Devices Sensors

IoT platforms

AWS IoT

STREAMS

Data streams

Stream storage

AWS re:Invent

aws

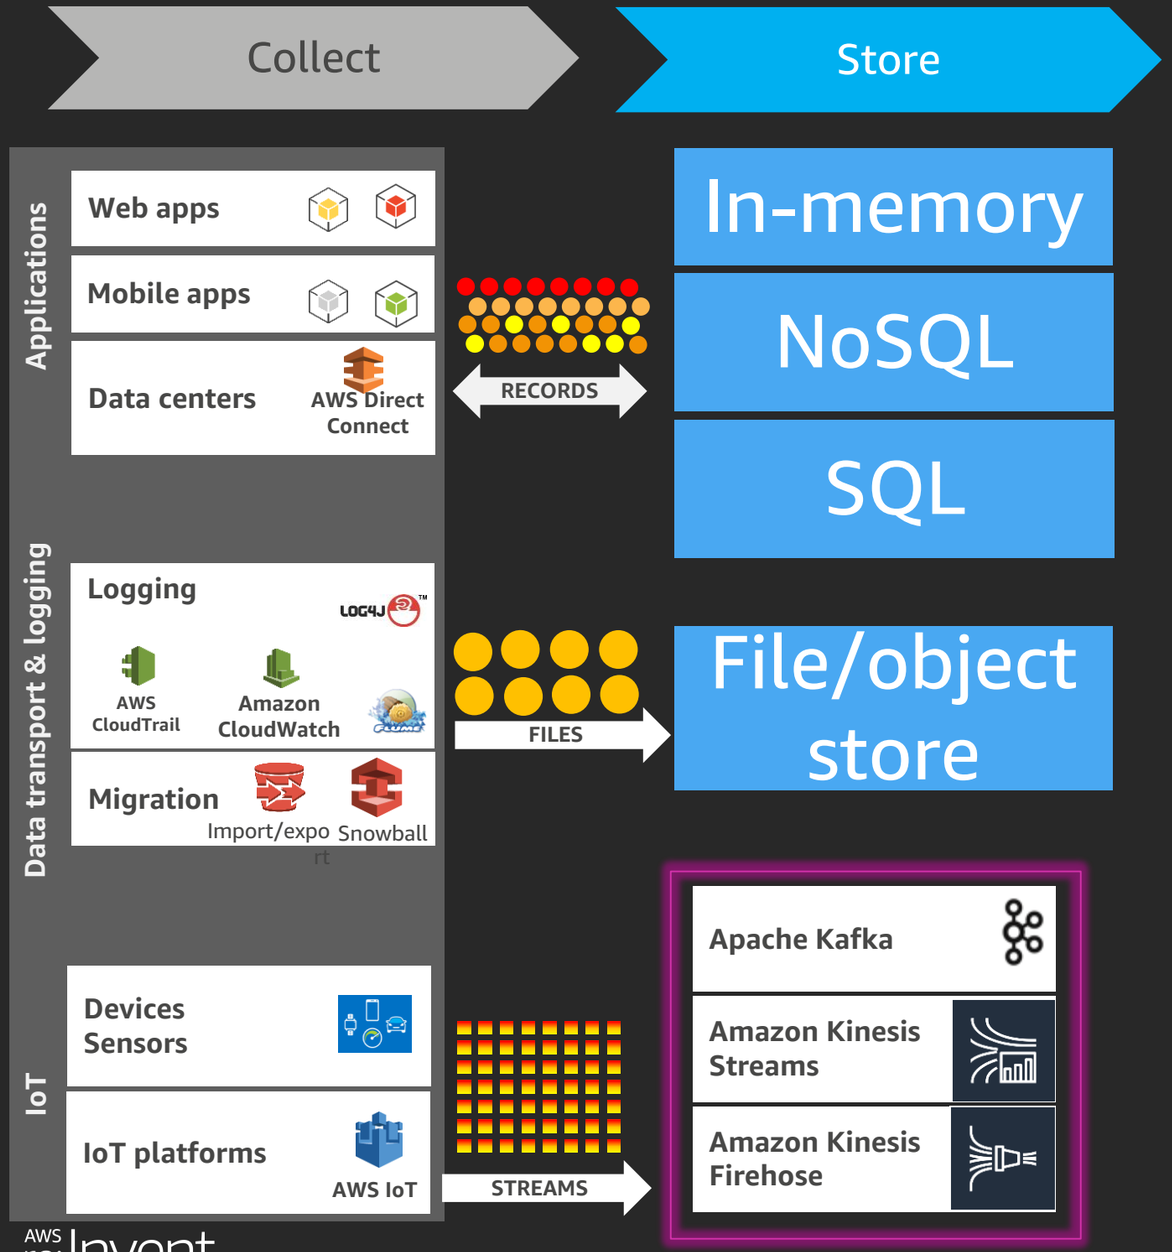

Stream storage

Apache Kafka

- High throughput distributed streaming platform

Amazon Kinesis Data Streams

- Managed stream storage

Amazon Kinesis Data Firehose

- Managed data delivery

aws

Which stream/message storage should I use?

	Amazon Kinesis Data Streams	Amazon Kinesis Data Firehose	Apache Kafka (on Amazon Elastic Compute Cloud [Amazon EC2])	Amazon Simple Queue Service (Amazon SQS) (Standard)	Amazon SQS (FIFO)
AWS managed	Yes	Yes	No	Yes	Yes
Guaranteed ordering	Yes	No	Yes	No	Yes
Delivery (deduping)	At least once	At least once	At least/At most/exactly once	At least once	Exactly once
Data retention period	7 days	N/A	Configurable	14 days	14 days
Availability	Three AZ	Three AZ	Configurable	Three AZ	Three AZ
Scale/throughput	No limit / ~ shards	No limit / automatic	No limit / ~ nodes	No limits / automatic	300 TPS / queue
Multiple consumers	Yes	No	Yes	No	No
Row/object size	1 MB	Destination row/object size	Configurable	256 KB	256 KB
Cost	Low	Low	Low (+admin)	Low-medium	Low-medium

File/object storage

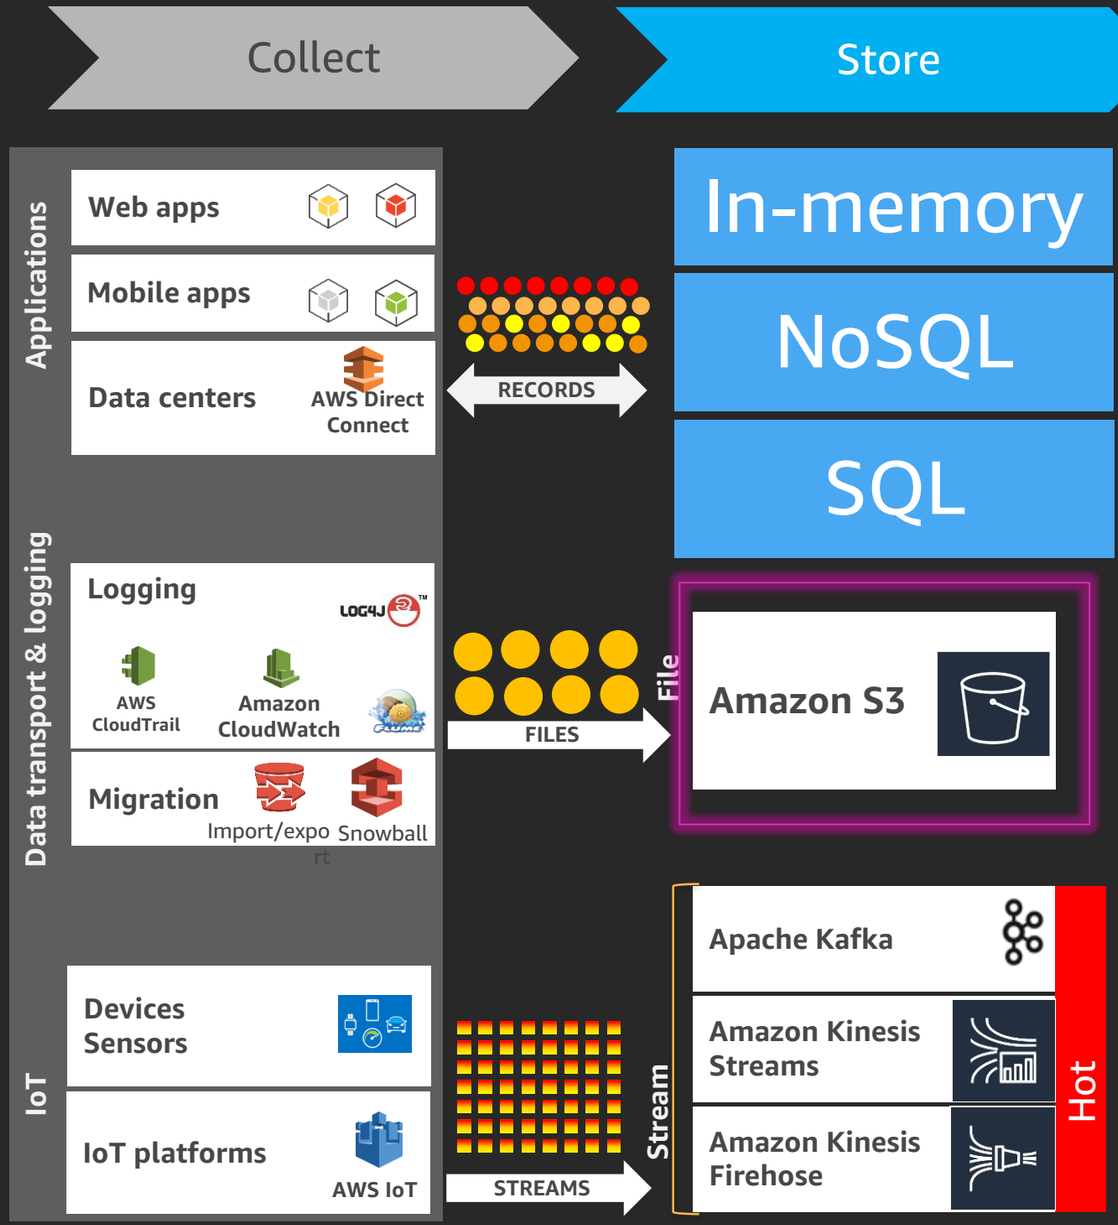

Amazon Simple Storage Service (Amazon S3)

- Managed object storage service built to store and retrieve any amount of data

aws

Use Amazon S3 as the storage for your data lake

- Natively supported by big data frameworks (Spark, Hive, Presto, and others)
- Decouple storage and compute
    - No need to run compute clusters for storage (unlike HDFS)
    - Can run transient Amazon EMR clusters with Amazon EC2 Spot Instances
    - Multiple & heterogeneous analysis clusters and services can use the same data
- Designed for 99.999999999% durability
- No need to pay for data replication within a region
- Secure – SSL, client/server-side encryption at rest
- Low cost

What about HDFS & data tiering?

- Use HDFS/local for working data sets
  - (for example, iterative read on the same data sets)
- Use Amazon S3 Standard for frequently accessed data
- Use Amazon S3 Standard–IA for less frequently accessed data
- Use Amazon Glacier for archiving cold data

Use S3 analytics to optimize tiering strategy

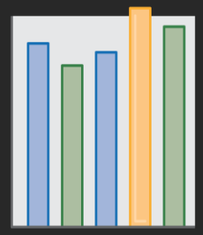

S3 analytics

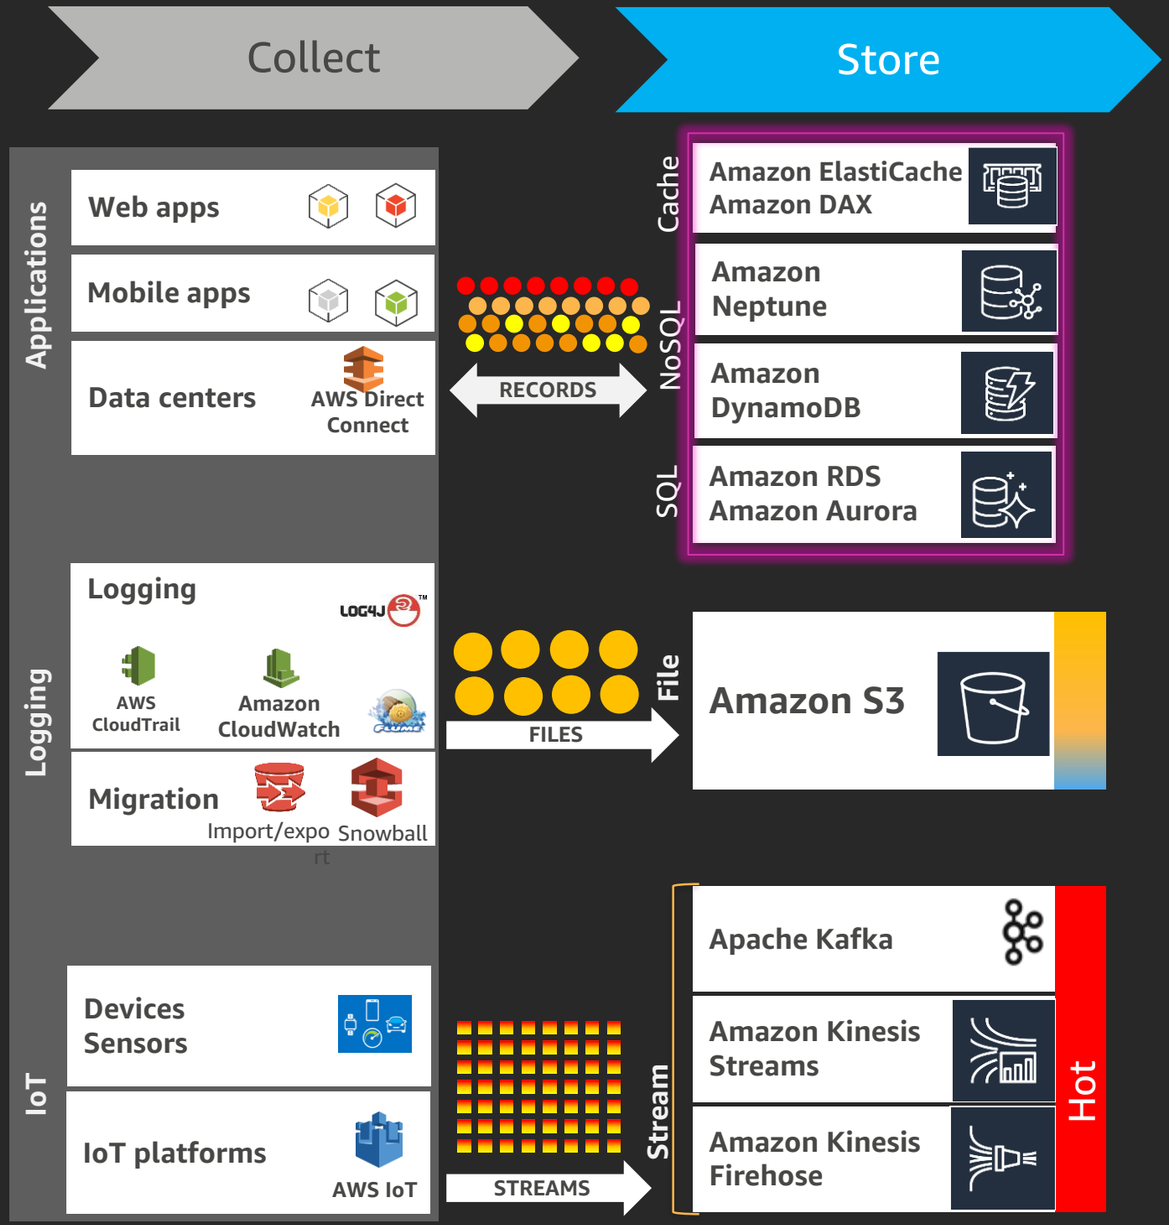

Cache & database

Amazon ElastiCache

- Managed Memcached or Redis service

Amazon DynamoDB Accelerator (DAX)

- Managed in-memory cache for DynamoDB

Amazon Neptune

- Managed graph DB

Amazon DynamoDB

- Managed key value/document DB

Amazon Relational Database Service (Amazon RDS)

- Managed relational database service

Anti-pattern

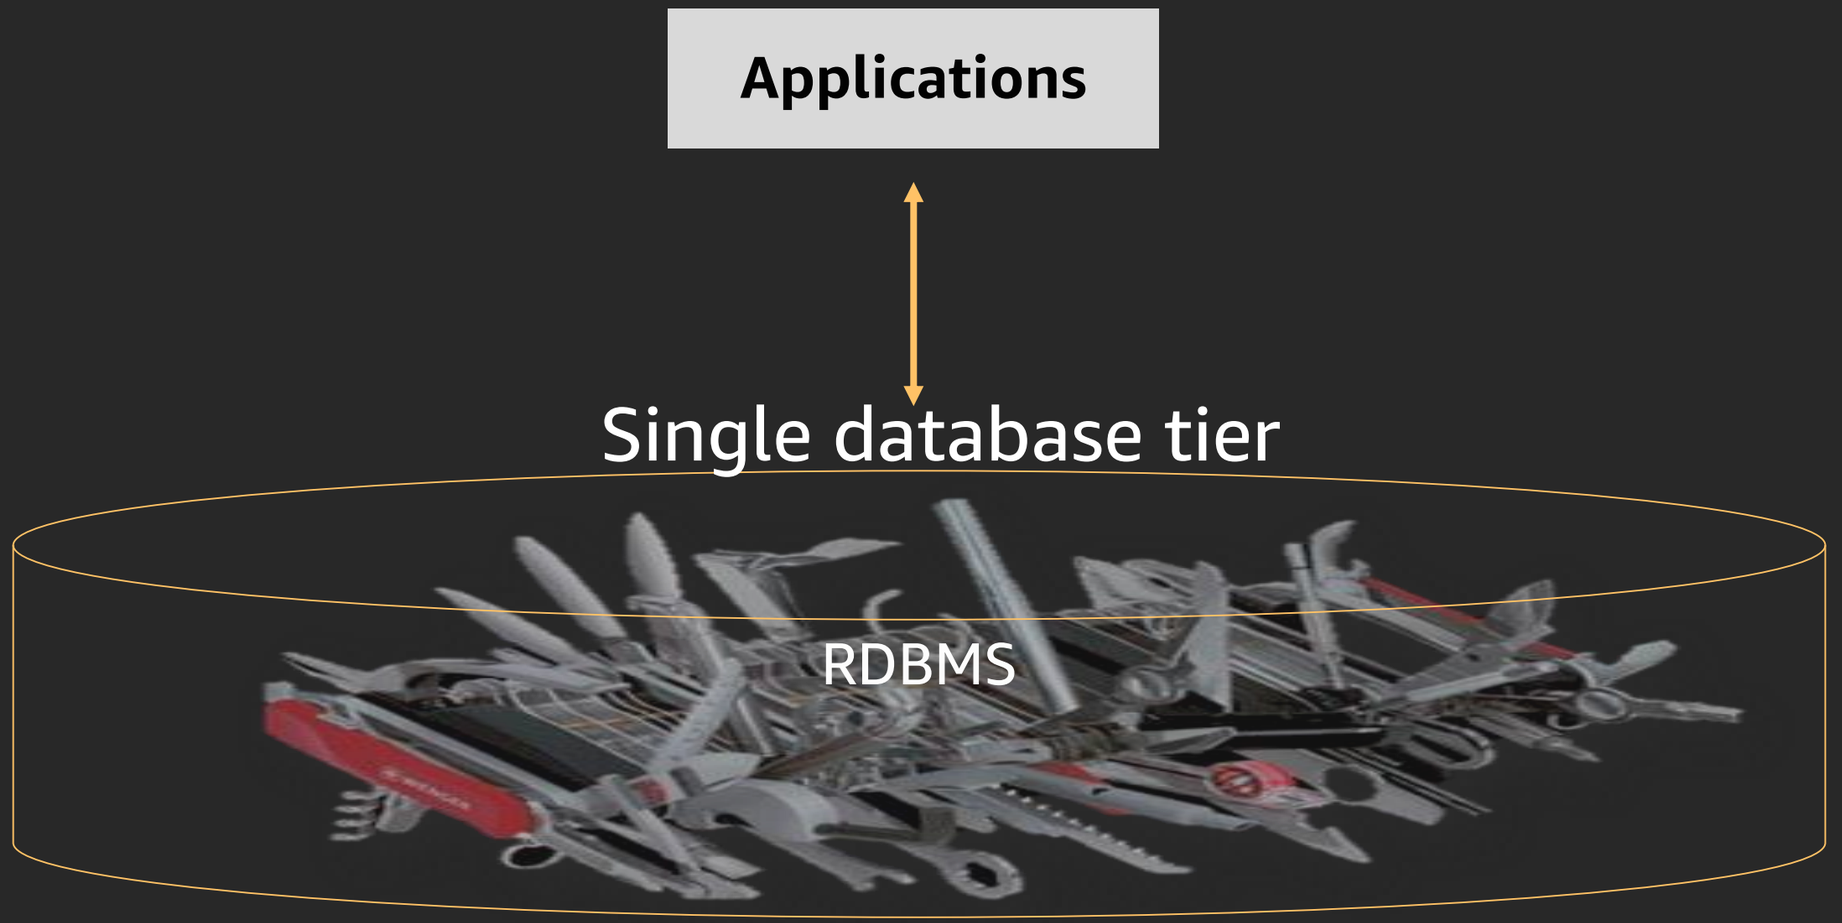

Best practice: Use the right tool for the job

Applications

Data tier

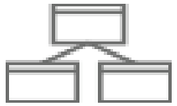

Relational

Referential integrity with strong consistency, transactions, and hardened scale

Complex query support via SQL

Key-value

Low-latency, key-based queries with high throughput and fast data ingestion

Simple query methods with filters

Document

Indexing and storing of documents with support for query on any property

Simple query with filters, projections and aggregates

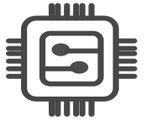

In-memory

Microsecond latency, key-based queries, specialized data structures

Simple query methods with filters

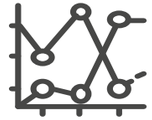

Graph

Creating and navigating relations between data easily and quickly

Easily express queries in terms of relations

Which data store should I use?

Ask yourself some questions

- What is the data structure?
- How will the data be accessed?
- What is the temperature of the data?
- What will the solution cost?

What is the data structure?

Data structure	What to use?
Fixed schema	SQL, NoSQL
Schema-free (JSON)	NoSQL, search
Key/value	In-memory, NoSQL
Graph	GraphDB

How will the data be accessed?

Access patterns	What to use?
Put/get (key, value)	In-memory, NoSQL
Simple relationships → 1:N, M:N	NoSQL
Multi-table joins, transaction, SQL	SQL
Faceting, search	Search
Graph traversal	GraphDB

What is the temperature of the data?

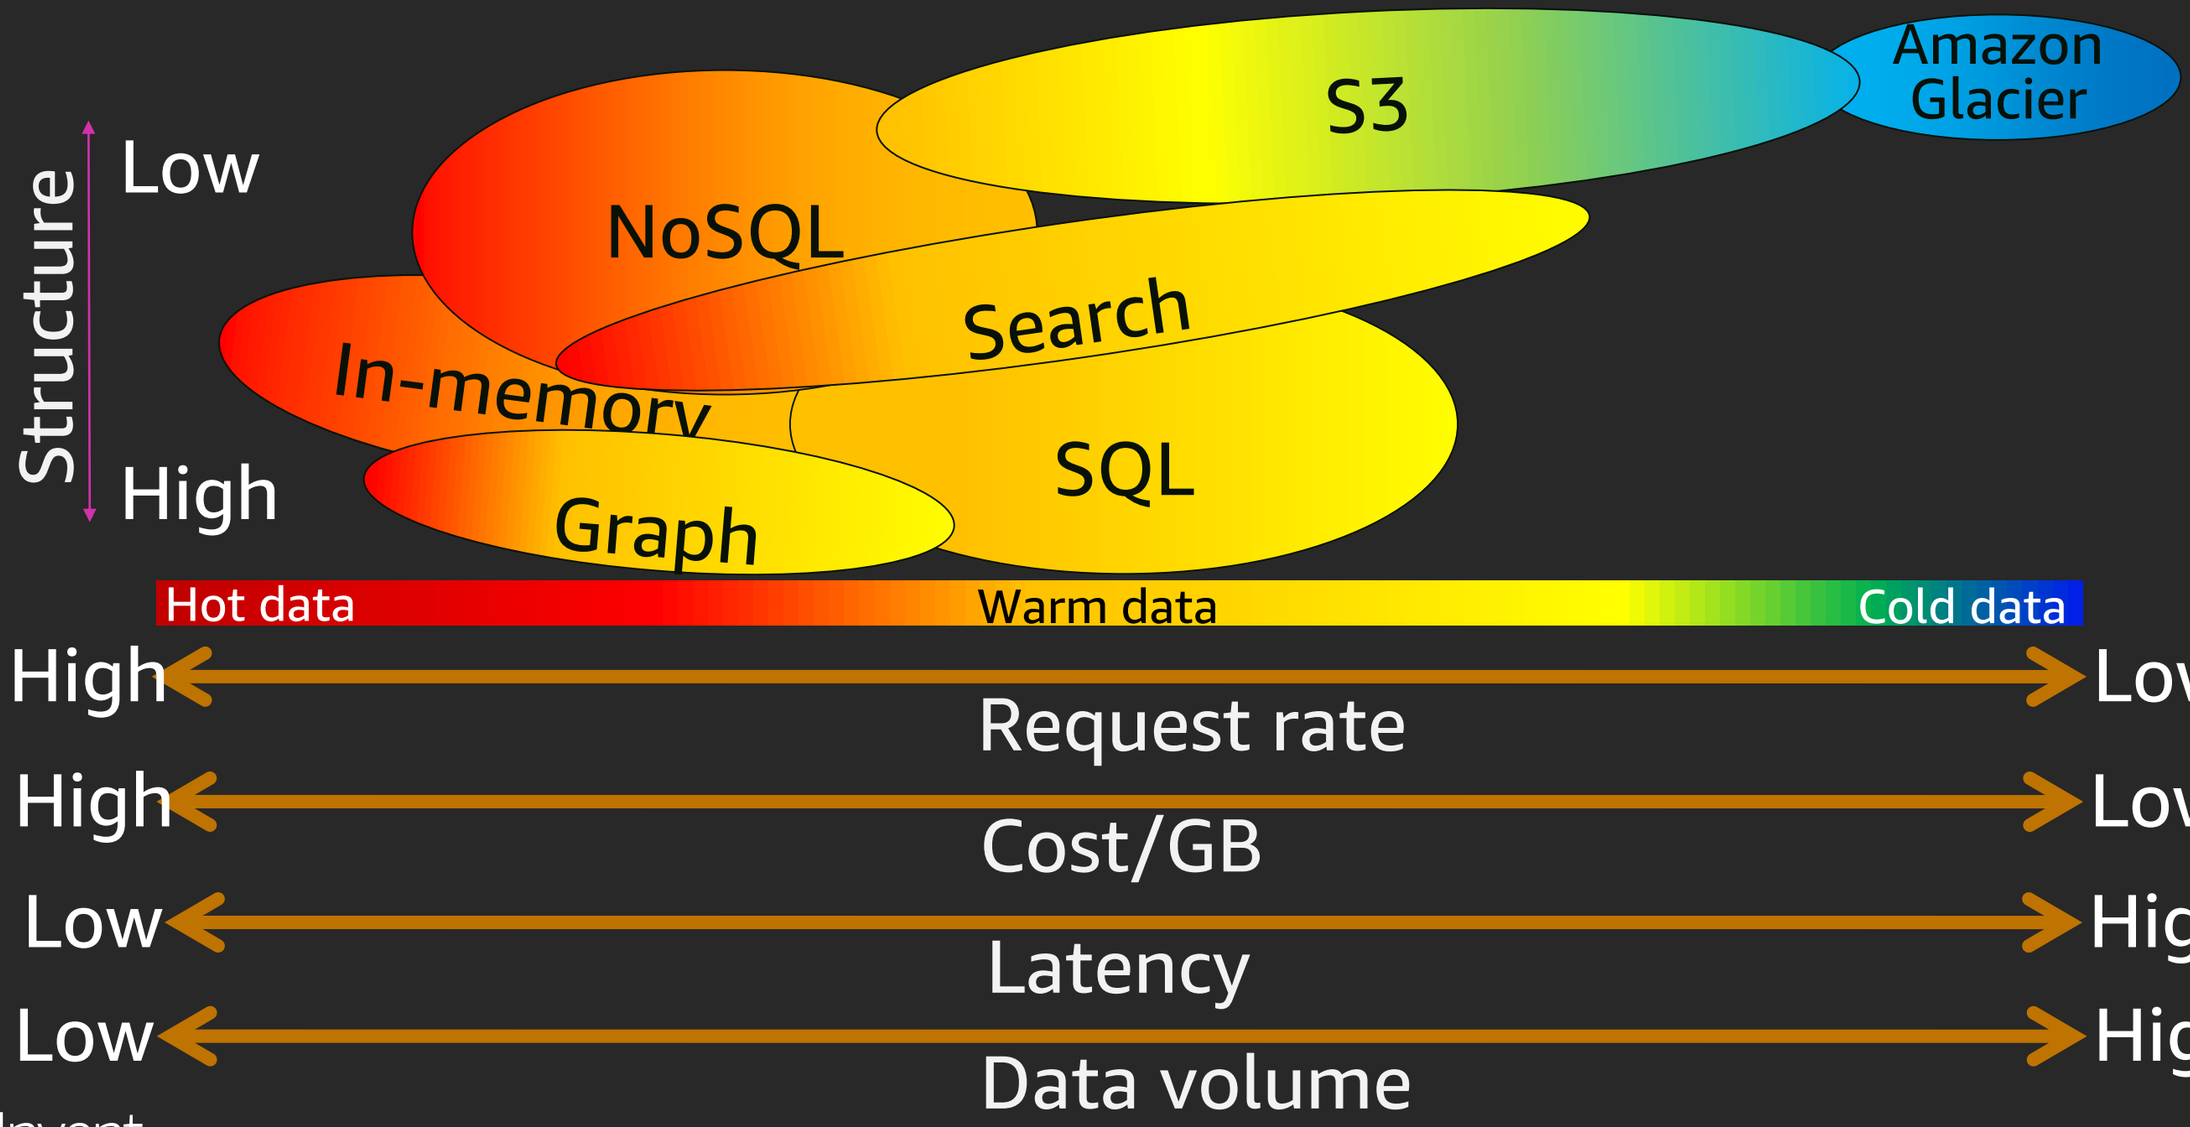

Database characteristics

	Amazon ElastiCache	Amazon DynamoDB + DAX	Amazon Aurora	Amazon RDS	Amazon Elasticsearch Service	Amazon Neptune	Amazon S3 + Amazon Glacier
Use cases	In memory caching	K/V lookups, document store	OLTP, transactional	OLTP, transactional	Log analysis, reverse indexing	Graph	File store
Performance	Ultra high request rate, Ultra low latency	Ultra high request rate, ultra low to low latency	Very high request rate, low latency	High request rate, low latency	Medium request rate, low latency	Medium request rate, low latency	High throughput
Shape	K/V	K/V and document	Relational	Relational	Documents	Node/edges	Files
Size	GB	TB, PB (no limits)	GB, mid TB	GB, low TB	GB, TB	GB, mid TB	GB, TB, PB, EB (no limits)
Cost / GB	$$	¢¢ - $$	¢¢	¢¢	¢¢	¢¢	¢- ¢4/10
Availability	2 AZ	3 AZ	3 AZ	2 AZ	1-2 AZ	3 AZ	3 AZ
VPC support	Inside VPC	VPC endpoint	Inside VPC	Inside VPC	Outside or inside VPC	Inside VPC	VPC endpoint

Hot data — Warm data — Cold data

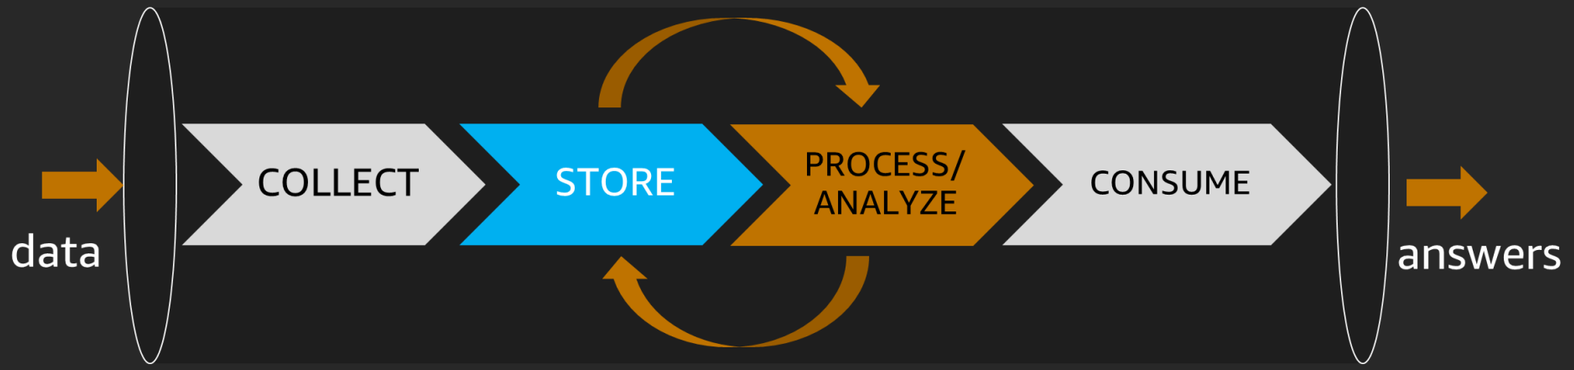

Process/ analyze

Interactive & batch analytics

Process/analyze

Amazon Elasticsearch Service

- Managed service for ElasticSearch

Amazon Redshift & Amazon Redshift Spectrum

- Managed data warehouse
- Spectrum enables querying S3

Amazon Athena

- Serverless interactive query service

Amazon EMR

- Managed Hadoop framework for running Apache Spark, Flink, Presto, Tez, Hive, Pig, HBase, and others

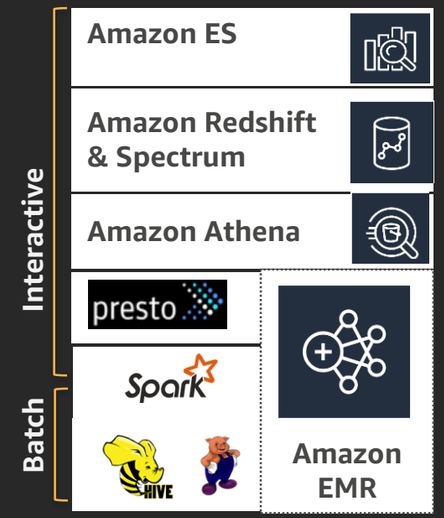

Stream/real-time analytics

Process/analyze

Spark Streaming on Amazon EMR

Amazon Kinesis Data Analytics

- Managed service for running SQL on streaming data

Amazon KCL

- Amazon Kinesis Client Library

AWS Lambda

- Run code serverless (without provisioning or managing servers)
- Services such as S3 can publish events to AWS Lambda
- AWS Lambda can pool event from a Kinesis

aws

Predictive analytics

Process/analyze

Predictive

Amazon ML

APPLICATION SERVICES

REKOGNITION | REKOGNITION VIDEO | POLLY | TRANSCRIBE | TRANSLATE | COMPREHEND | LEX

PLATFORMS

Amazon SageMaker | Amazon Mechanical Turk | Amazon Deep Learning AMIs

FRAMEWORKS & INFRASTRUCTURE

Caffe2 | CNTK | mxnet | PYTORCH | TensorFlow | KERAS | GLUON

P3 — NVIDIA Tesla V100 GPU accelerated for AI/ML training

C5 — Compute intensive instances for AI/ML inference

Machine Learning AMIs

Greengrass ML

Developers

Data scientists
Deep learning experts

aws

Which analytics should I use?

Batch

Takes minutes to hours

Example: Daily/weekly/monthly reports

Amazon EMR (MapReduce, Hive, Pig, Spark)

Interactive

Takes seconds

Example: Self-service dashboards

Amazon Redshift, Amazon Athena, Amazon EMR (Presto, Spark)

Stream

Takes milliseconds to seconds

Example: Fraud alerts, one-minute metrics

Amazon EMR (Spark Streaming), Amazon Kinesis Data Analytics, Amazon KCL, AWS Lambda, and others

Predictive

Takes milliseconds (real-time) to minutes (batch)

Example: Fraud detection, forecasting demand, speech recognition

Amazon SageMaker, Amazon Polly, Amazon Rekognition, Amazon Transcribe, Amazon Translate, Amazon EMR (Spark ML), Amazon Deep Learning AMI (MXNet, TensorFlow, Theano, Torch, CNTK, and Caffe)

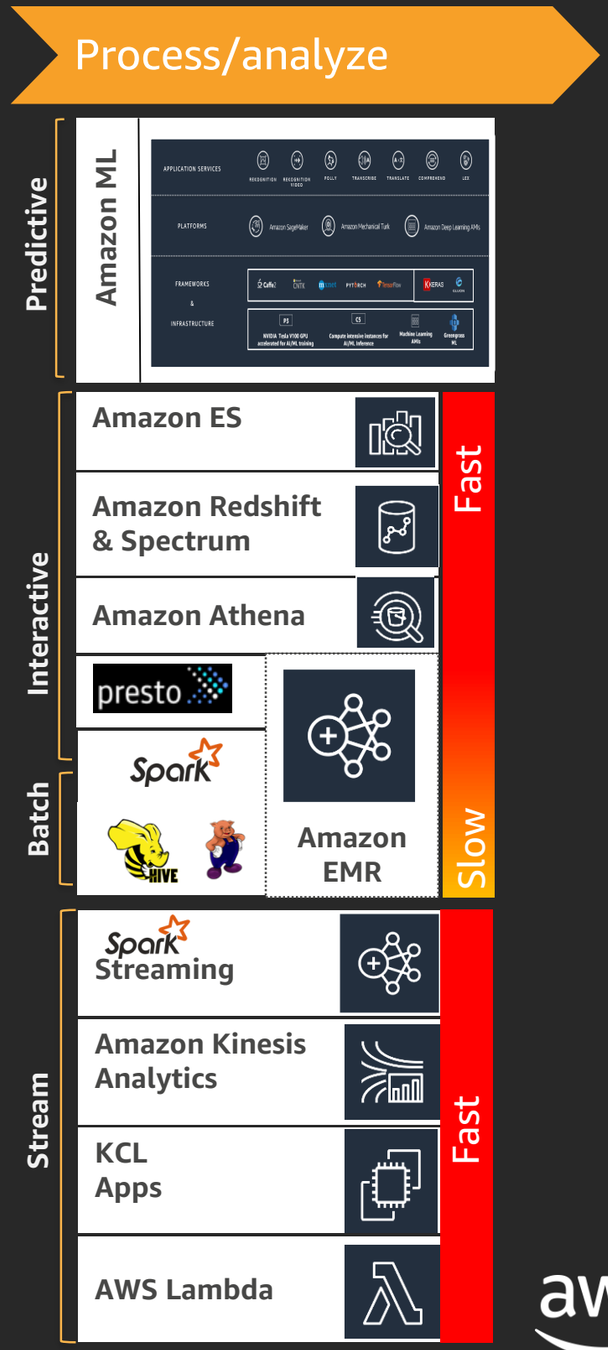

Which stream processing technology should I use?

	Amazon EMR (Spark Streaming)	KCL application	Amazon Kinesis analytics	AWS Lambda
Managed service	Yes	No (EC2 + Auto Scaling)	Yes	Yes
Serverless	No	No	Yes	Yes
Scale / throughput	No limits / ~ nodes	No limits / ~ nodes	No Limits / automatic	No limits / automatic
Availability	Single AZ	Multi-AZ	Multi-AZ	Multi-AZ
Programming languages	Java, Python, Scala	Java, others through MultiLangDaemon	ANSI SQL with extensions	Node.js, Java, Python, .Net Core
Sliding window functions	Build-in	App needs to implement	Built-in	No
Reliability	KCL and Spark checkpoints	Managed by Amazon KCL	Managed by Amazon Kinesis Data Analytics	Managed by AWS Lambda

	Amazon Redshift	Amazon Redshift Spectrum	Amazon Athena	Amazon EMR		
				Presto	Spark	Hive
Use case	Optimized for data warehousing	Query S3 data from Amazon Redshift	Interactive queries over S3 data	Interactive query	General purpose	Batch
Scale/throughput	~Nodes	~Nodes	Automatic	~ Nodes		
Managed service	Yes	Yes	Yes, Serverless	Yes		
Storage	Local storage	Amazon S3	Amazon S3	Amazon S3, HDFS		
Optimization	Columnar storage, data compression, and zone maps	AVRO, PARQUET TEXT, SEQ RCFILE, ORC, etc.	AVRO, PARQUET TEXT, SEQ RCFILE, ORC, etc.	Framework dependent		
Metadata	Amazon Redshift catalog	AWS Glue catalog	AWS Glue catalog	AWS Glue catalog or Hive Meta-store		
Auth/access controls	IAM, users, groups, and access controls	IAM, users, groups, and access controls	IAM	IAM, LDAP, & Kerberos		
UDF support	Yes (Scalar)	Yes (Scalar)	No	Yes		

Fastest — Fast — Slow

ELT/ETL: Preparing raw data for consumption

Raw data stored in data lake

Preparation

- *Normalized*
- *Partitioned*
- *Compressed*
- *Storage optimized*

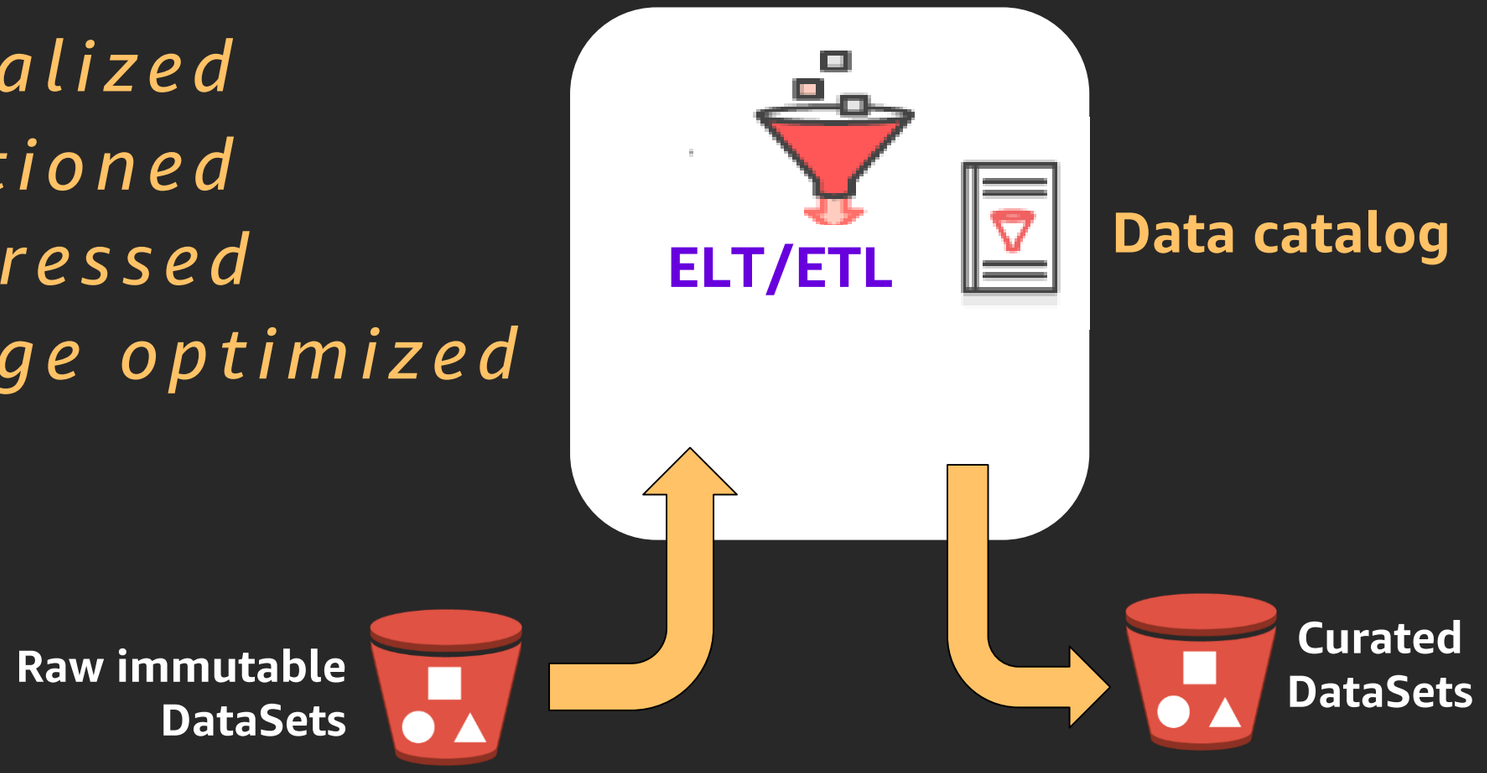

Data lake on AWS

Demonstration—Why ELT?

Which ELT/ETL tool should I use?

	AWS Glue ETL	AWS Data Pipeline	Amazon Data Migration Service (Amazon DMS)	Amazon EMR	Apache NiFi	Partner solution
Use case	Serverless ETL	Data workflow service	Migrate databases (and to/from data lake)	Customize developed Hadoop/Spark ETL	Automate the flow of data between systems	Rich partner ecosystem for ETL
Scale/throughput	~DPUs	~Nodes through EMR cluster	EC2 instance type	~Nodes	Self managed	Self manager or through partner
Managed service	Clusterless	Managed	Managed EC2 on your behalf	Managed EC2 on your behalf	Self managed on Amazon EMR or Marketplace	Self manager or through partner
Data sources	S3, RDBMS, Amazon RedShift, DynamoDB	S3, JDBC, DynamoDB, custom	RDBMS, data warehouses, Amazon S3* (*limited)	Various Hadoop/Spark managed	Various through rich processor framework	Various
Skills needed	Wizard for simple mapping, code snippets for advanced ETL	Wizard and code snippets	Wizard and drag/drop	Hadoop/Spark coding	NiFi processors and some coding	Self manager or through partner

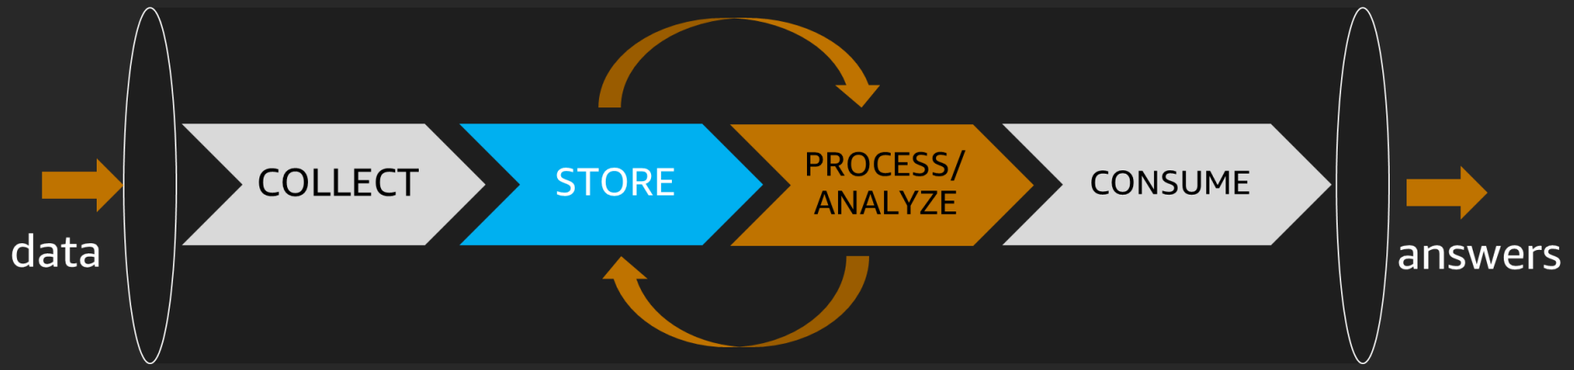

Consume

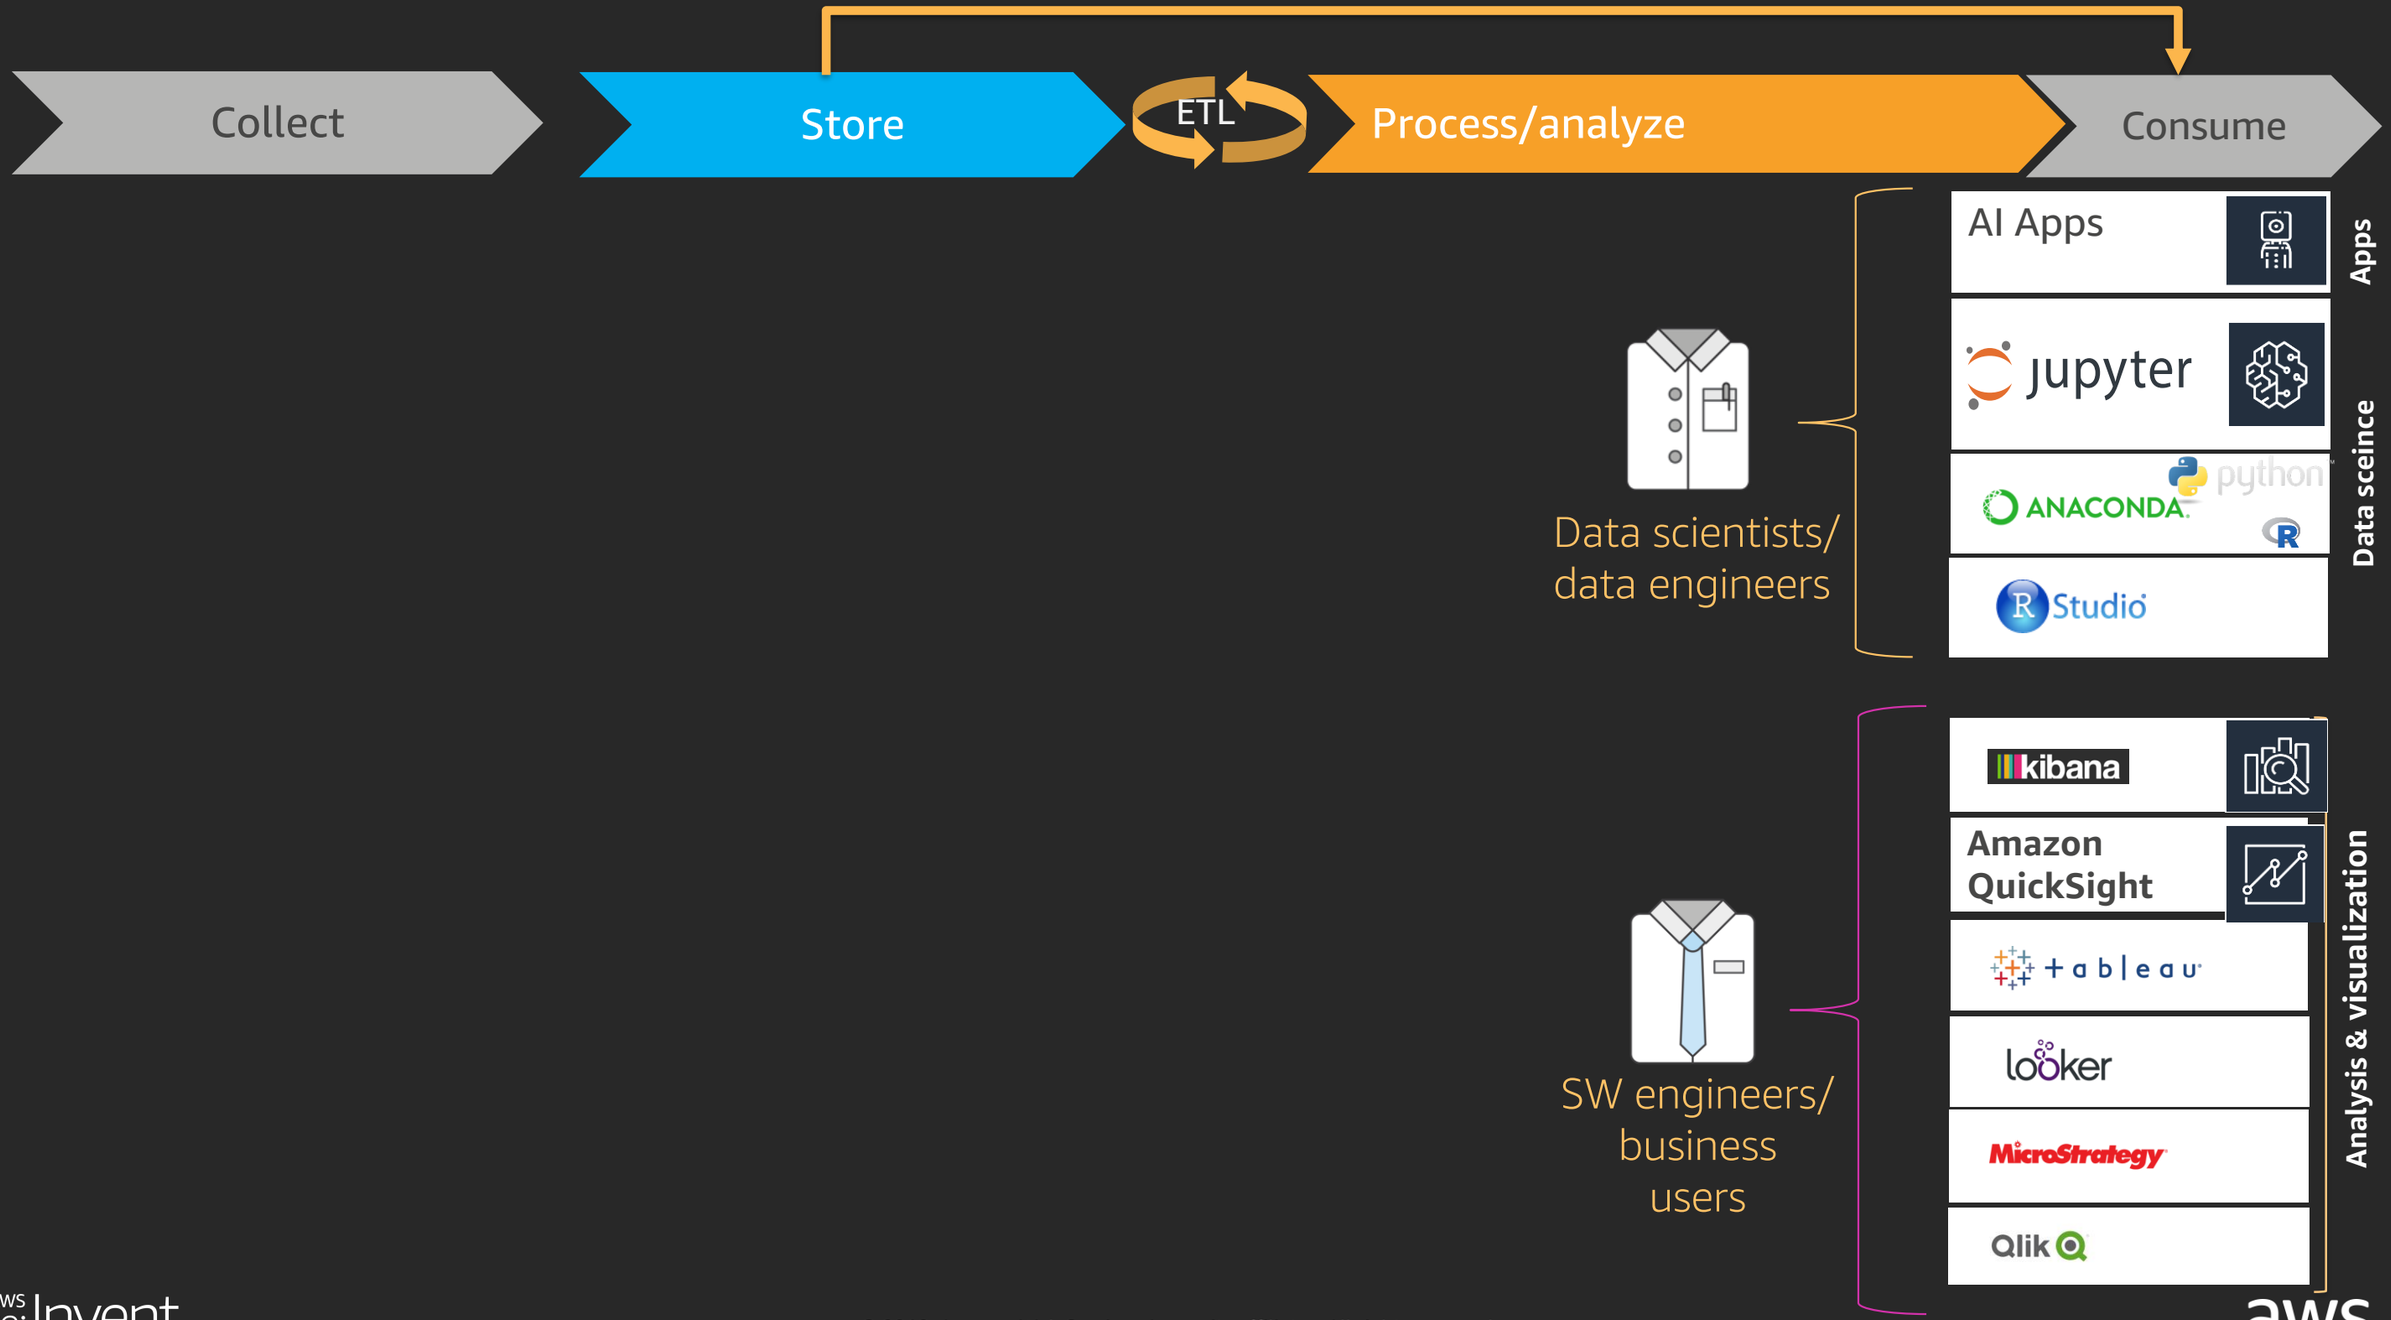

Putting it all together

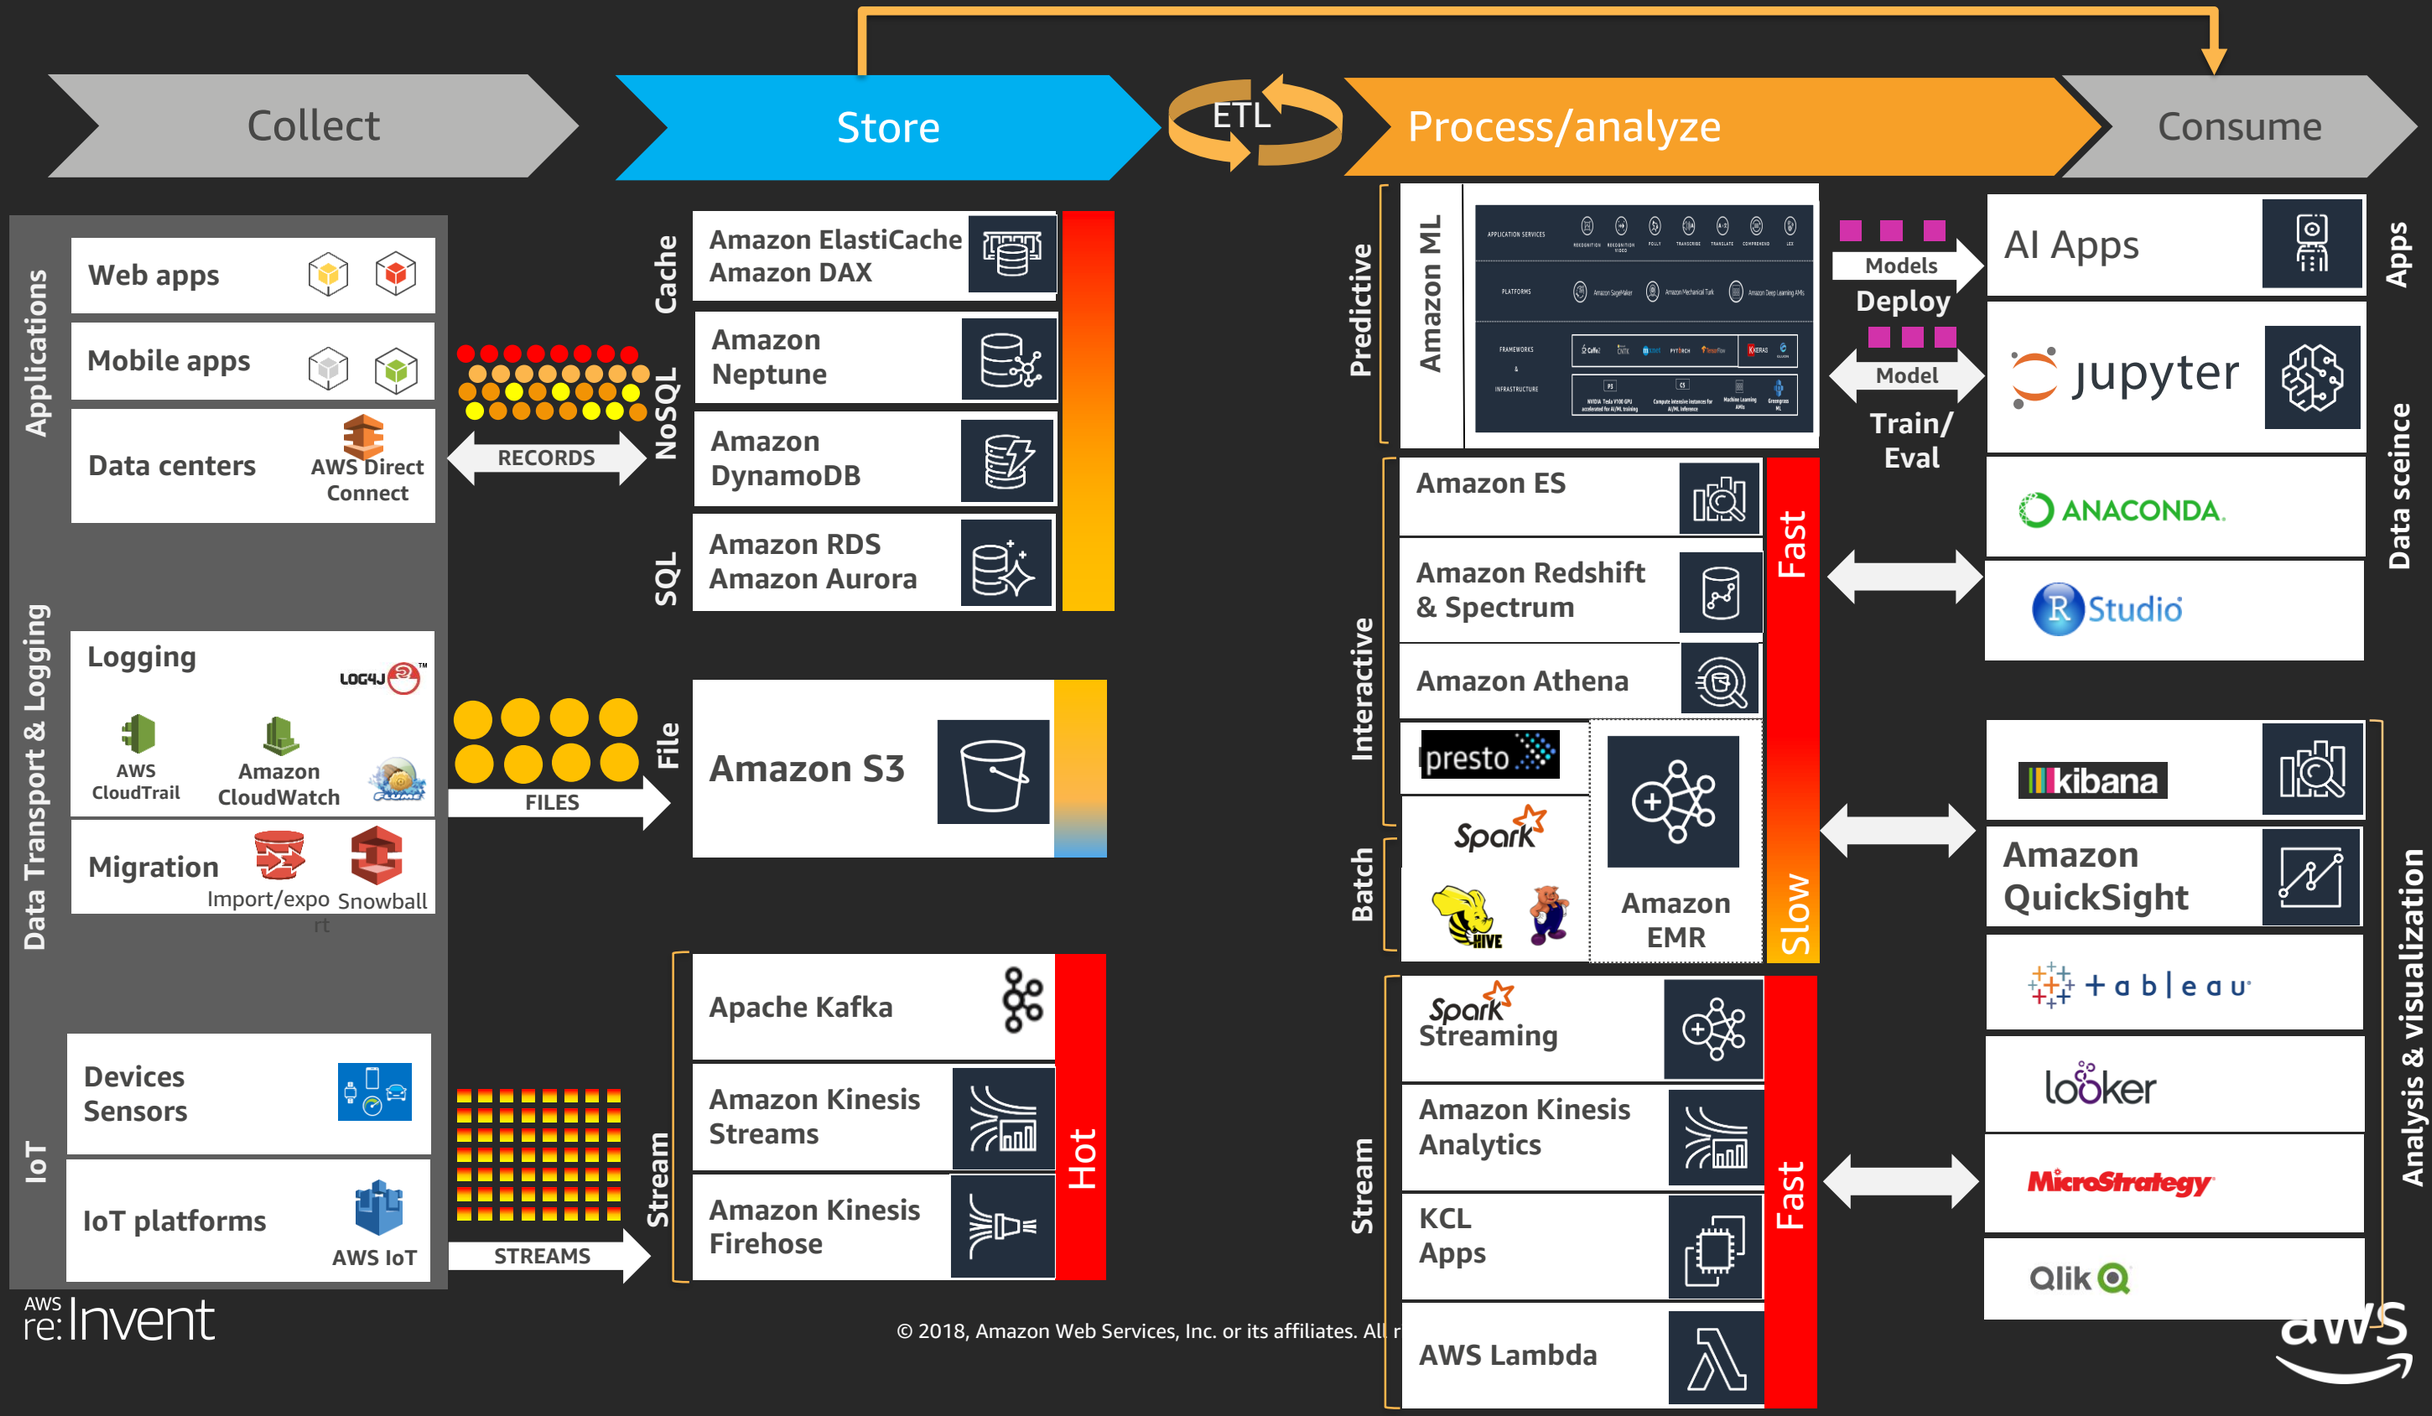
Collect
Store
ETL
Process/analyze
Consume
Applications
Web apps
Mobile apps
Data centers
AWS Direct Connect
Data Transport & Logging
Logging
LOG4J
AWS CloudTrail
Amazon CloudWatch
Migration
Import/export
Snowball
IoT
Devices Sensors
IoT platforms
AWS IoT
RECORDS
FILES
STREAMS
Cache
Amazon ElastiCache
Amazon DAX
NoSQL
Amazon Neptune
Amazon DynamoDB
SQL
Amazon RDS
Amazon Aurora
File
Amazon S3
Stream
Apache Kafka
Amazon Kinesis Streams
Amazon Kinesis Firehose
Hot
Predictive
Amazon ML
Models
Deploy
Model
Train/ Eval
Interactive
Amazon ES
Amazon Redshift & Spectrum
Amazon Athena
presto
Batch
Spark
HIVE
Amazon EMR
Fast
Slow
Stream
Spark Streaming
Amazon Kinesis Analytics
KCL Apps
AWS Lambda
Fast
Apps
AI Apps
jupyter
Data sceince
ANACONDA.
RStudio
Analysis & visualization
kibana
Amazon QuickSight
tableau
looker
MicroStrategy
Qlik
aws
AWS re:Invent
© 2018, Amazon Web Services, Inc. or its affiliates. All r

Design patterns

AWS re:Invent

aws

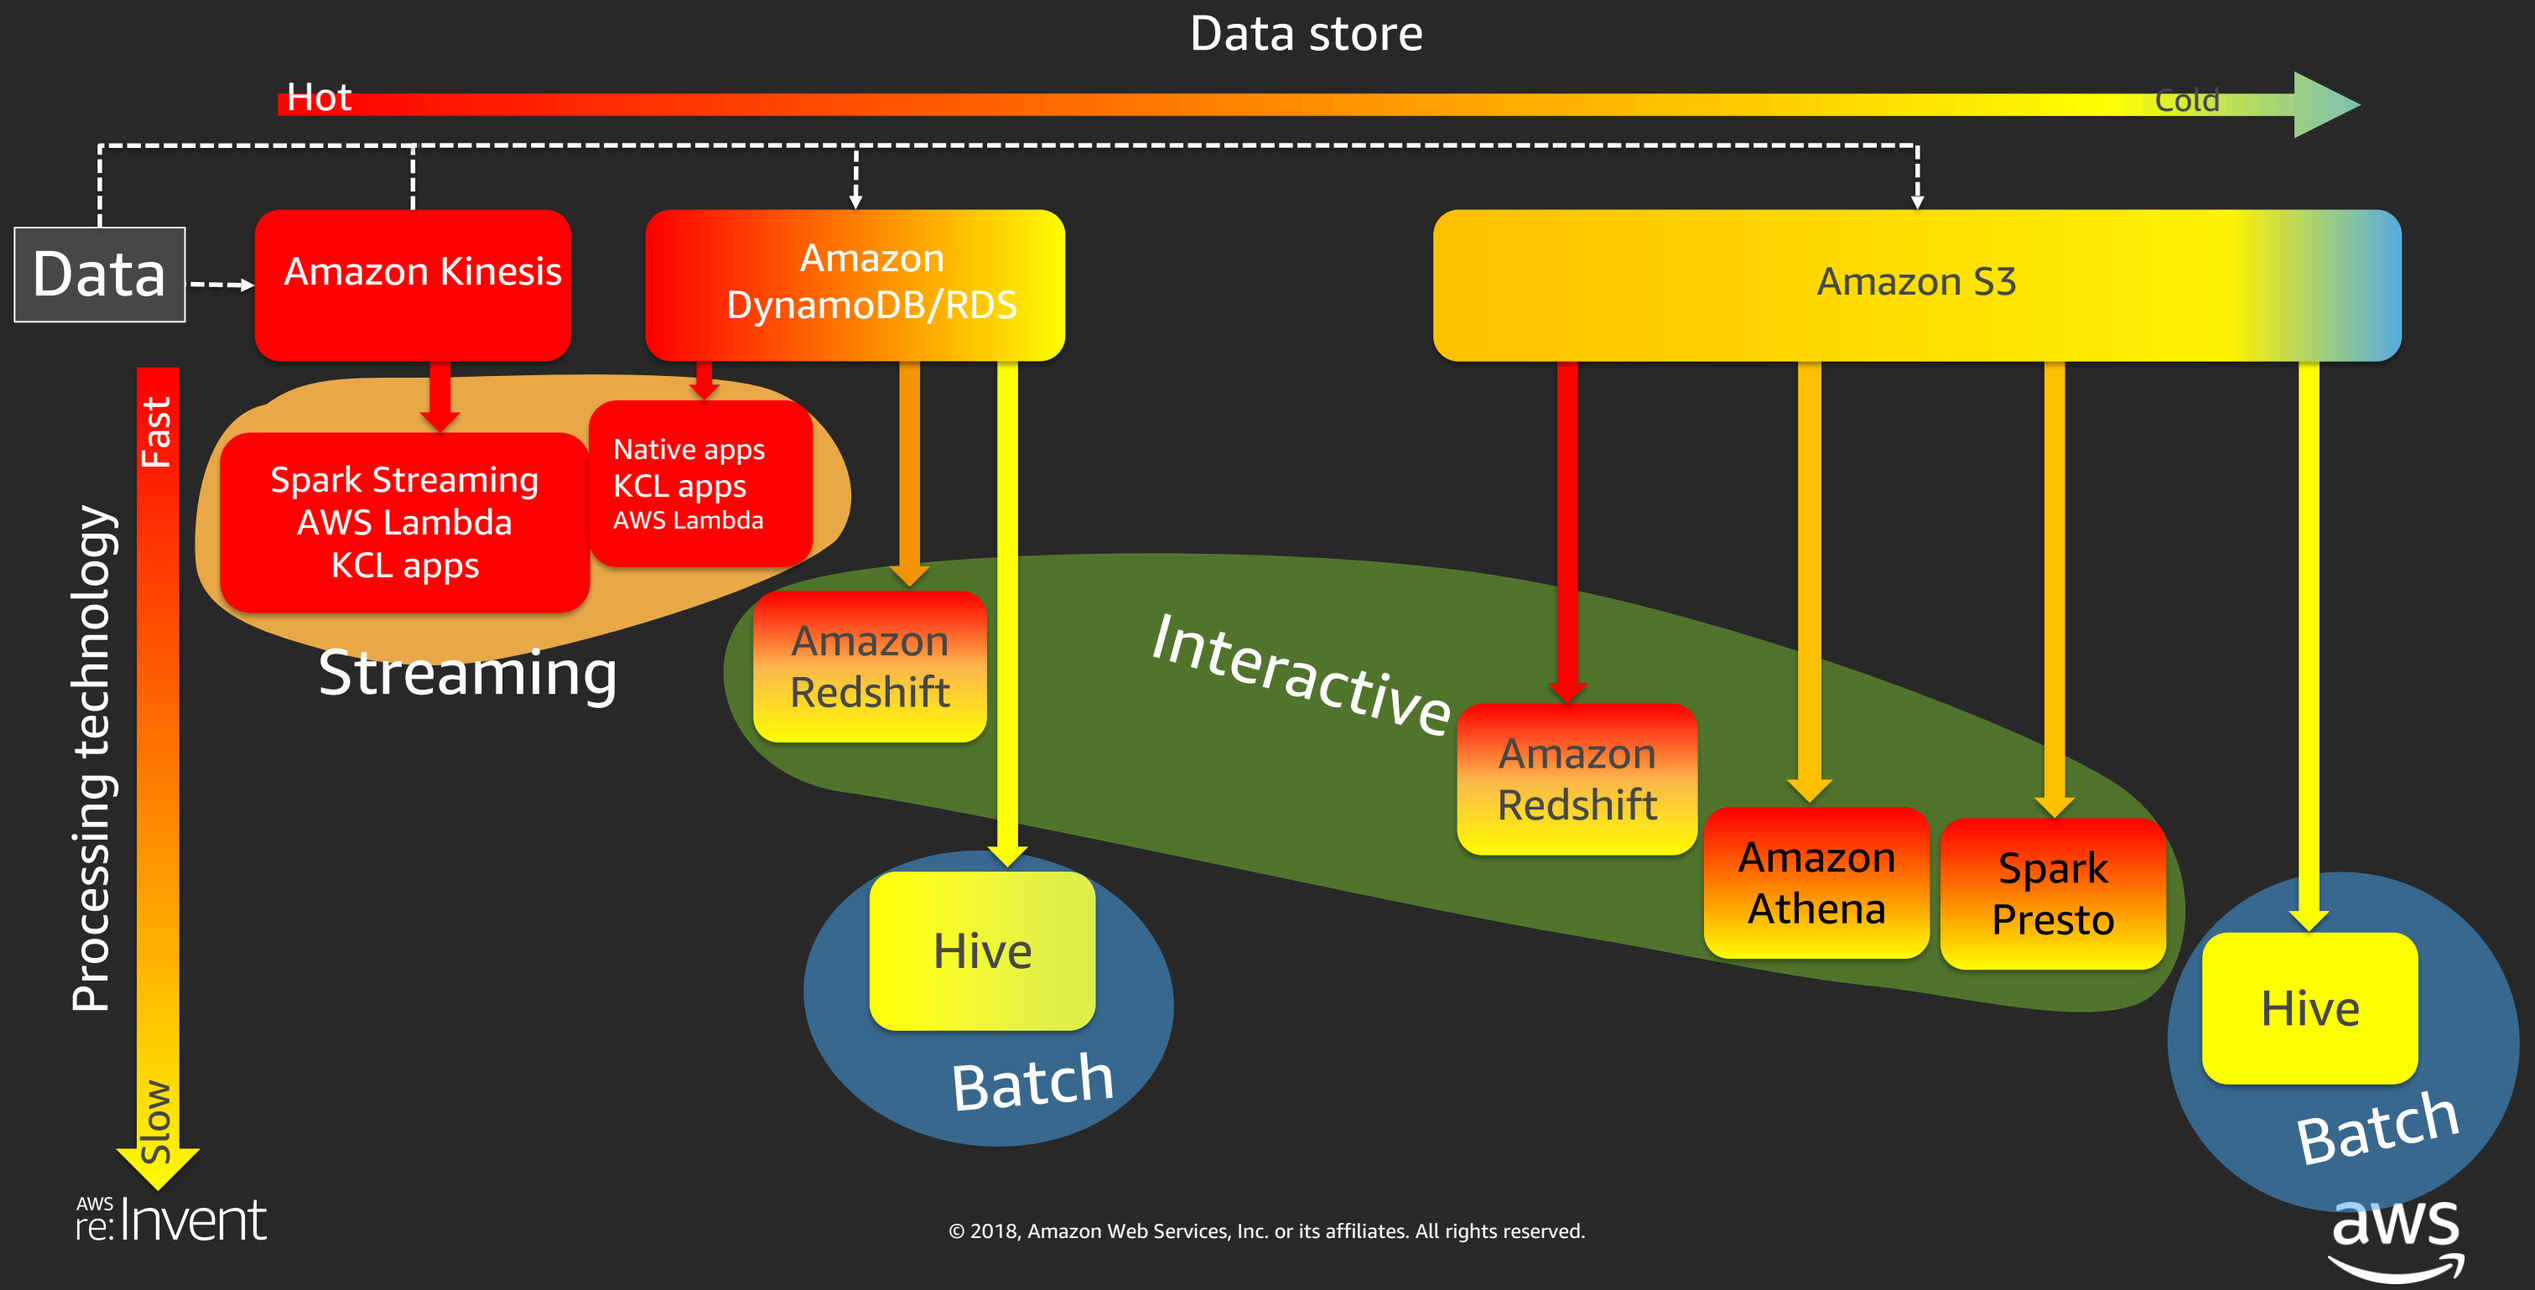
Data store
Hot
Cold
Data
Amazon Kinesis
Amazon DynamoDB/RDS
Amazon S3
Fast
Processing technology
Slow
Spark Streaming
AWS Lambda
KCL apps
Native apps
KCL apps
AWS Lambda
Streaming
Amazon Redshift
Interactive
Amazon Redshift
Amazon Athena
Spark Presto
Hive
Batch
Hive
Batch
AWS re:Invent
© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.
aws

Streaming analytics

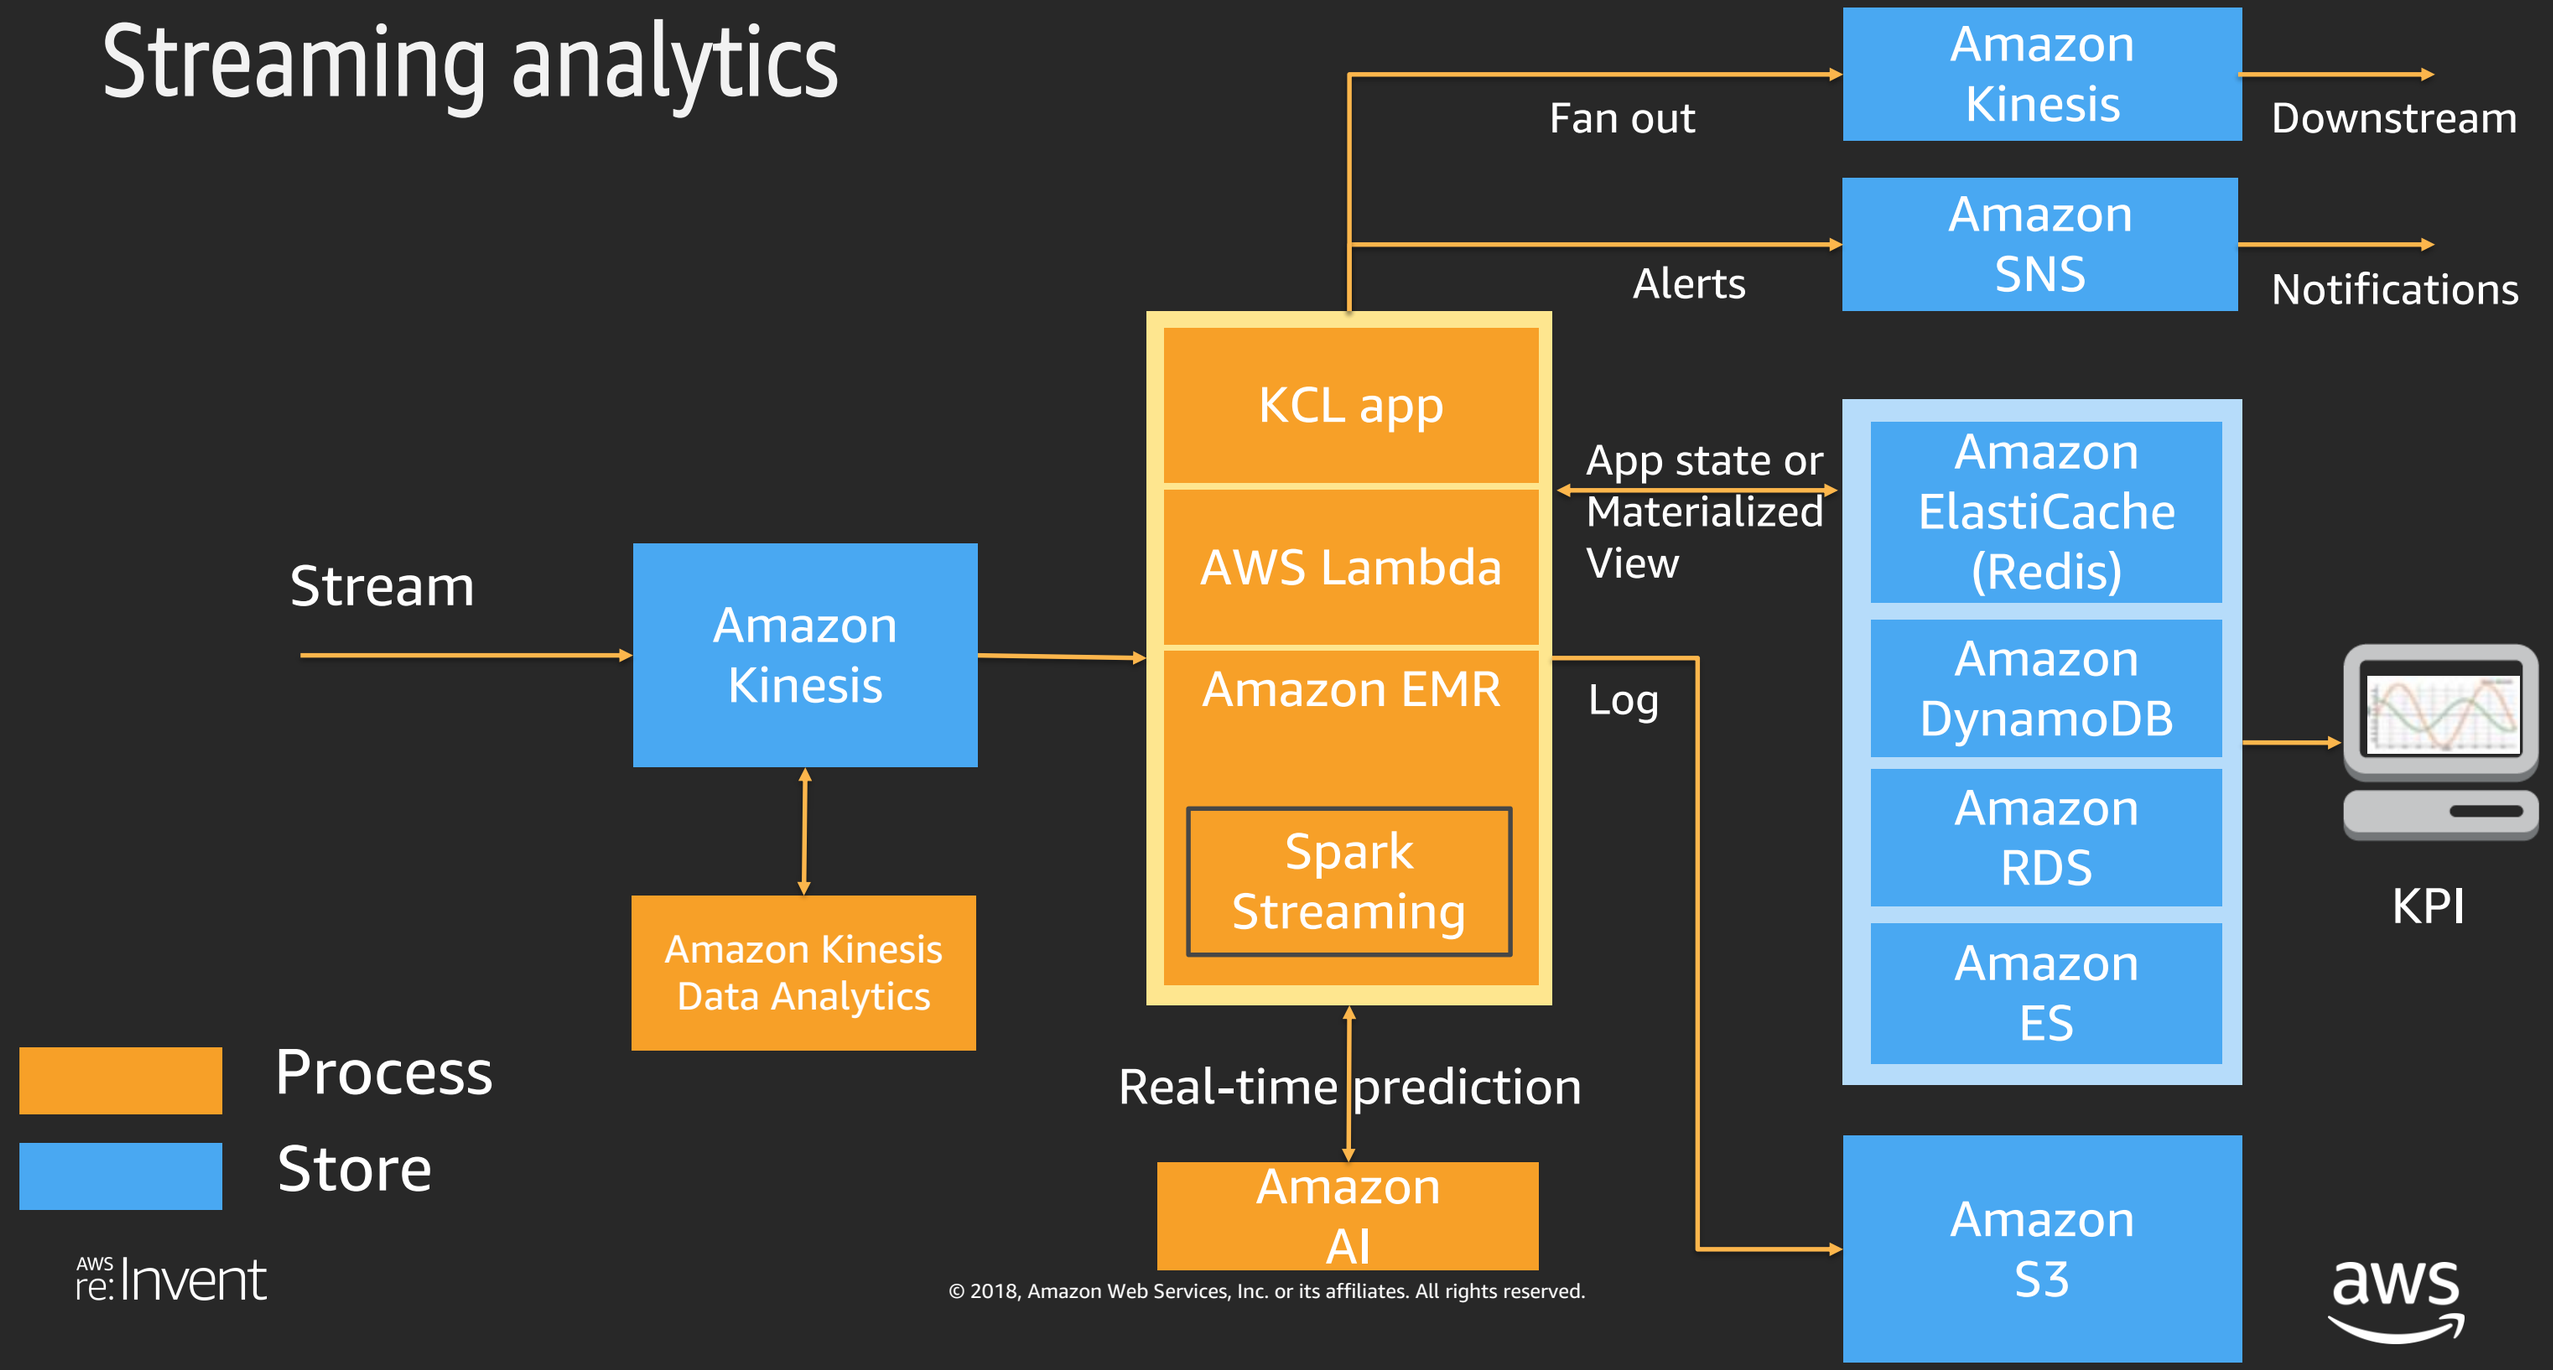

Hearst's serverless data pipeline

HEARST

Media and information company

Television

ESPN A&E THE HISTORY CHANNEL LIFETIME

Magazines

COSMOPOLITAN CAR and DRIVER BAZAAR Bicycling

News papers

Chron San Francisco Chronicle

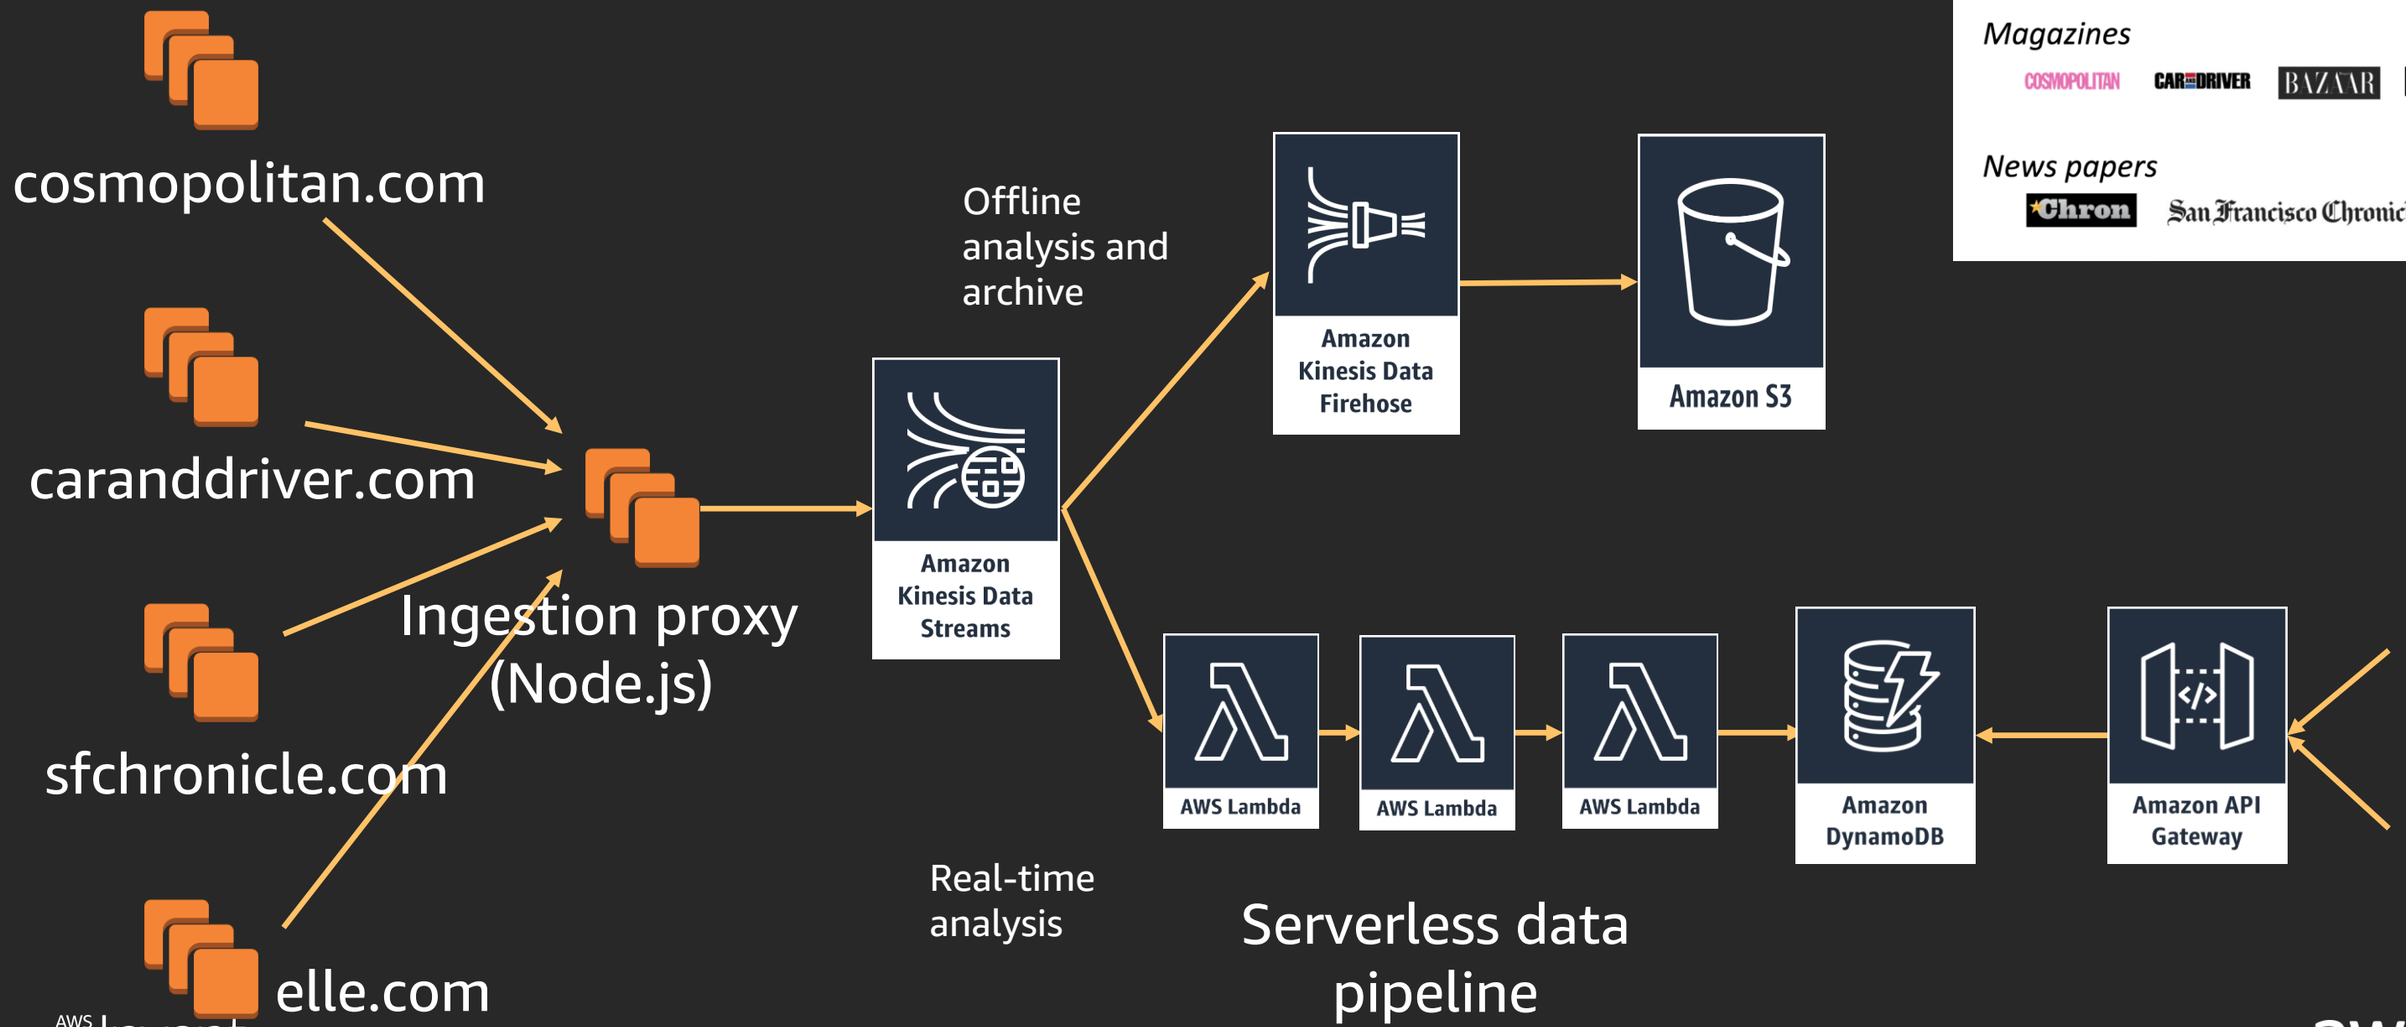

Yieldmo's data ingestion and real-time analysis architecture

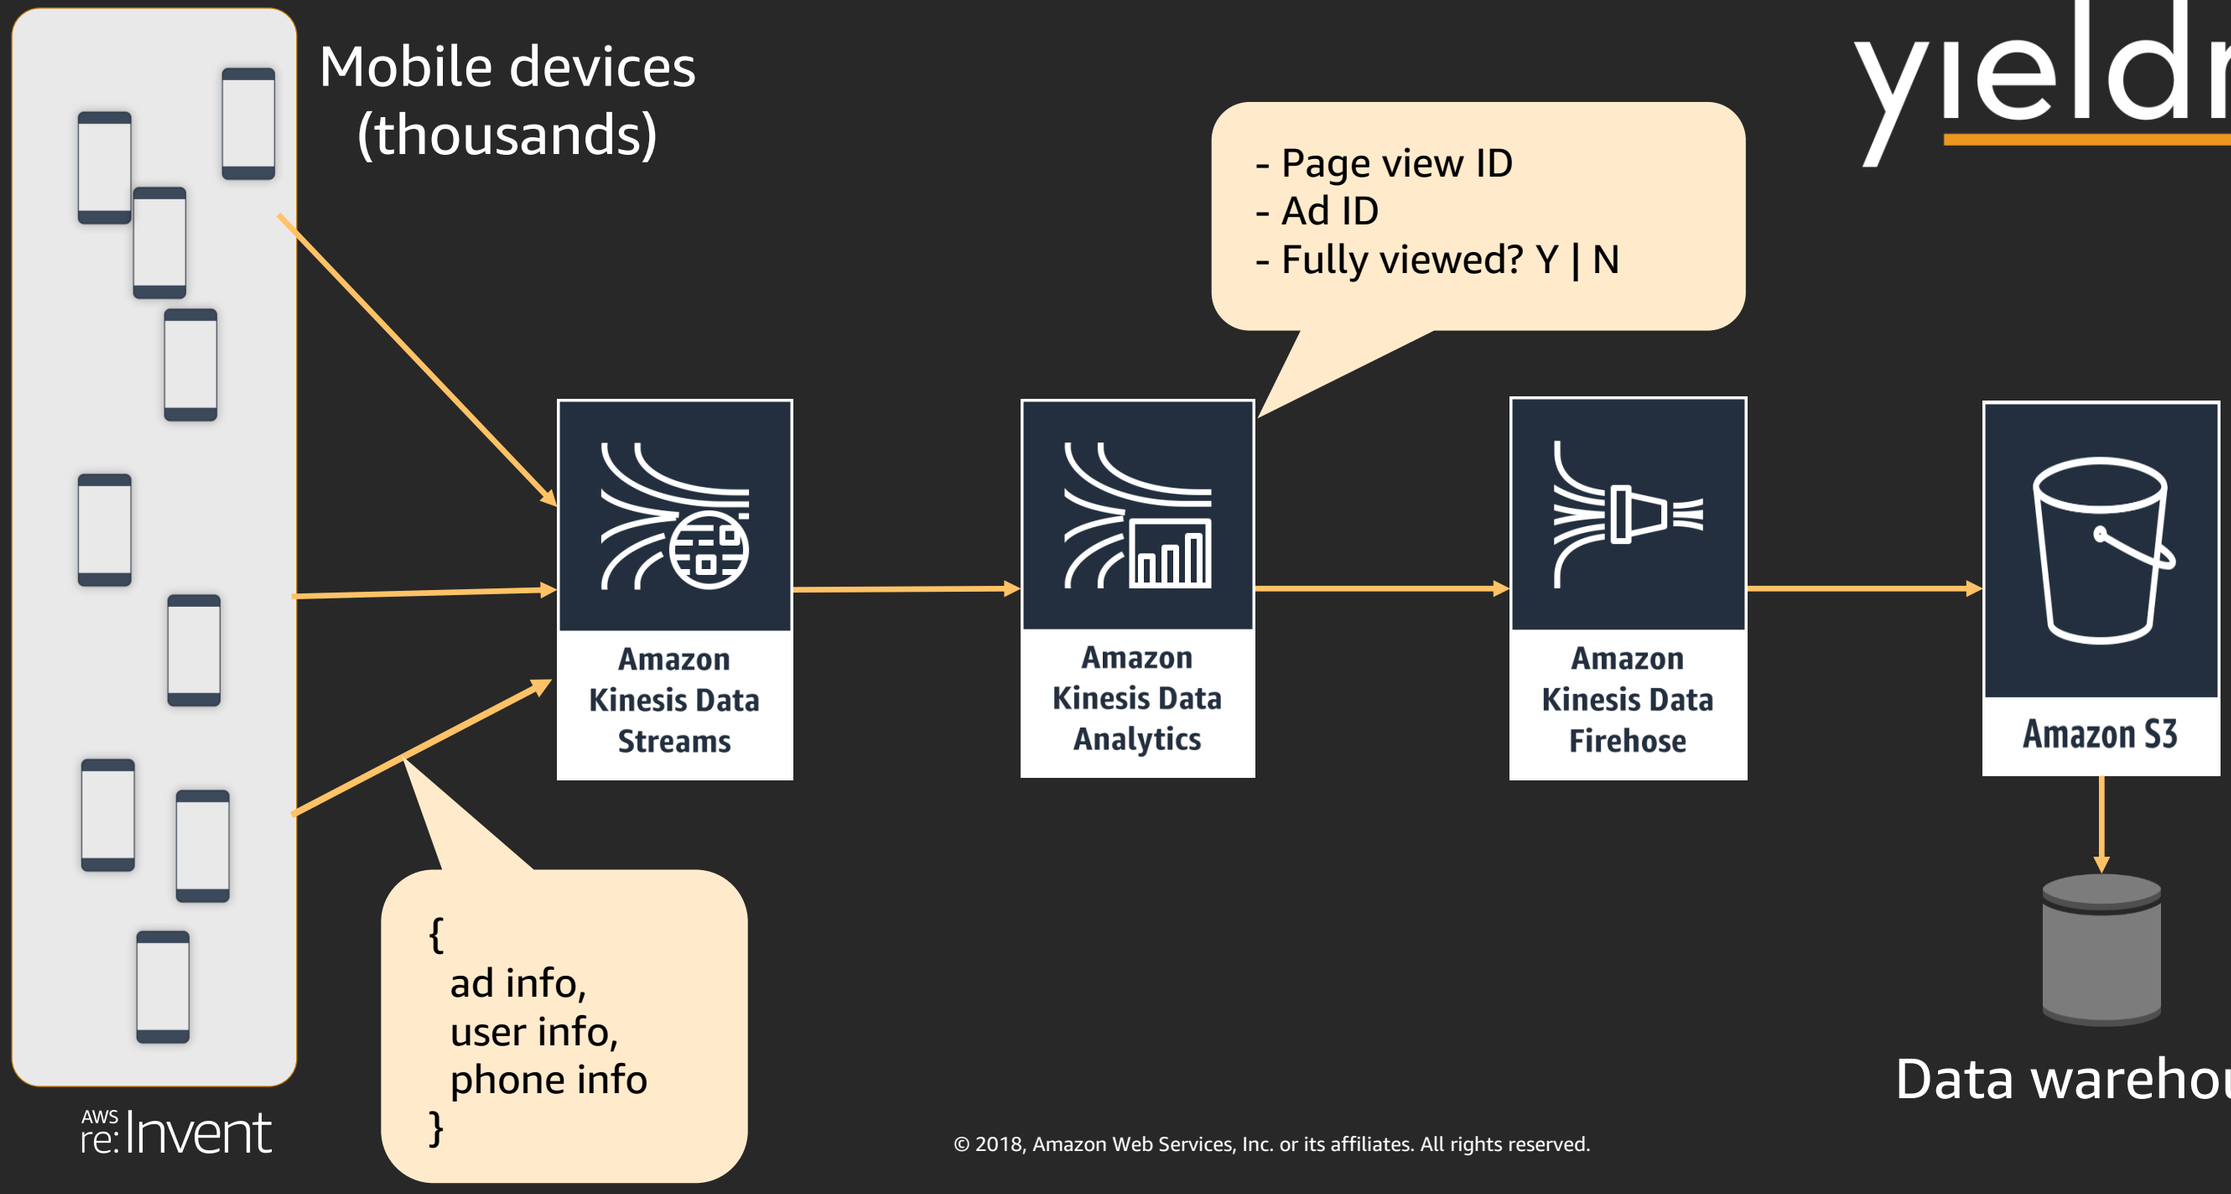

yieldmo

AWS re:Invent

aws

Interactive & batch analytics

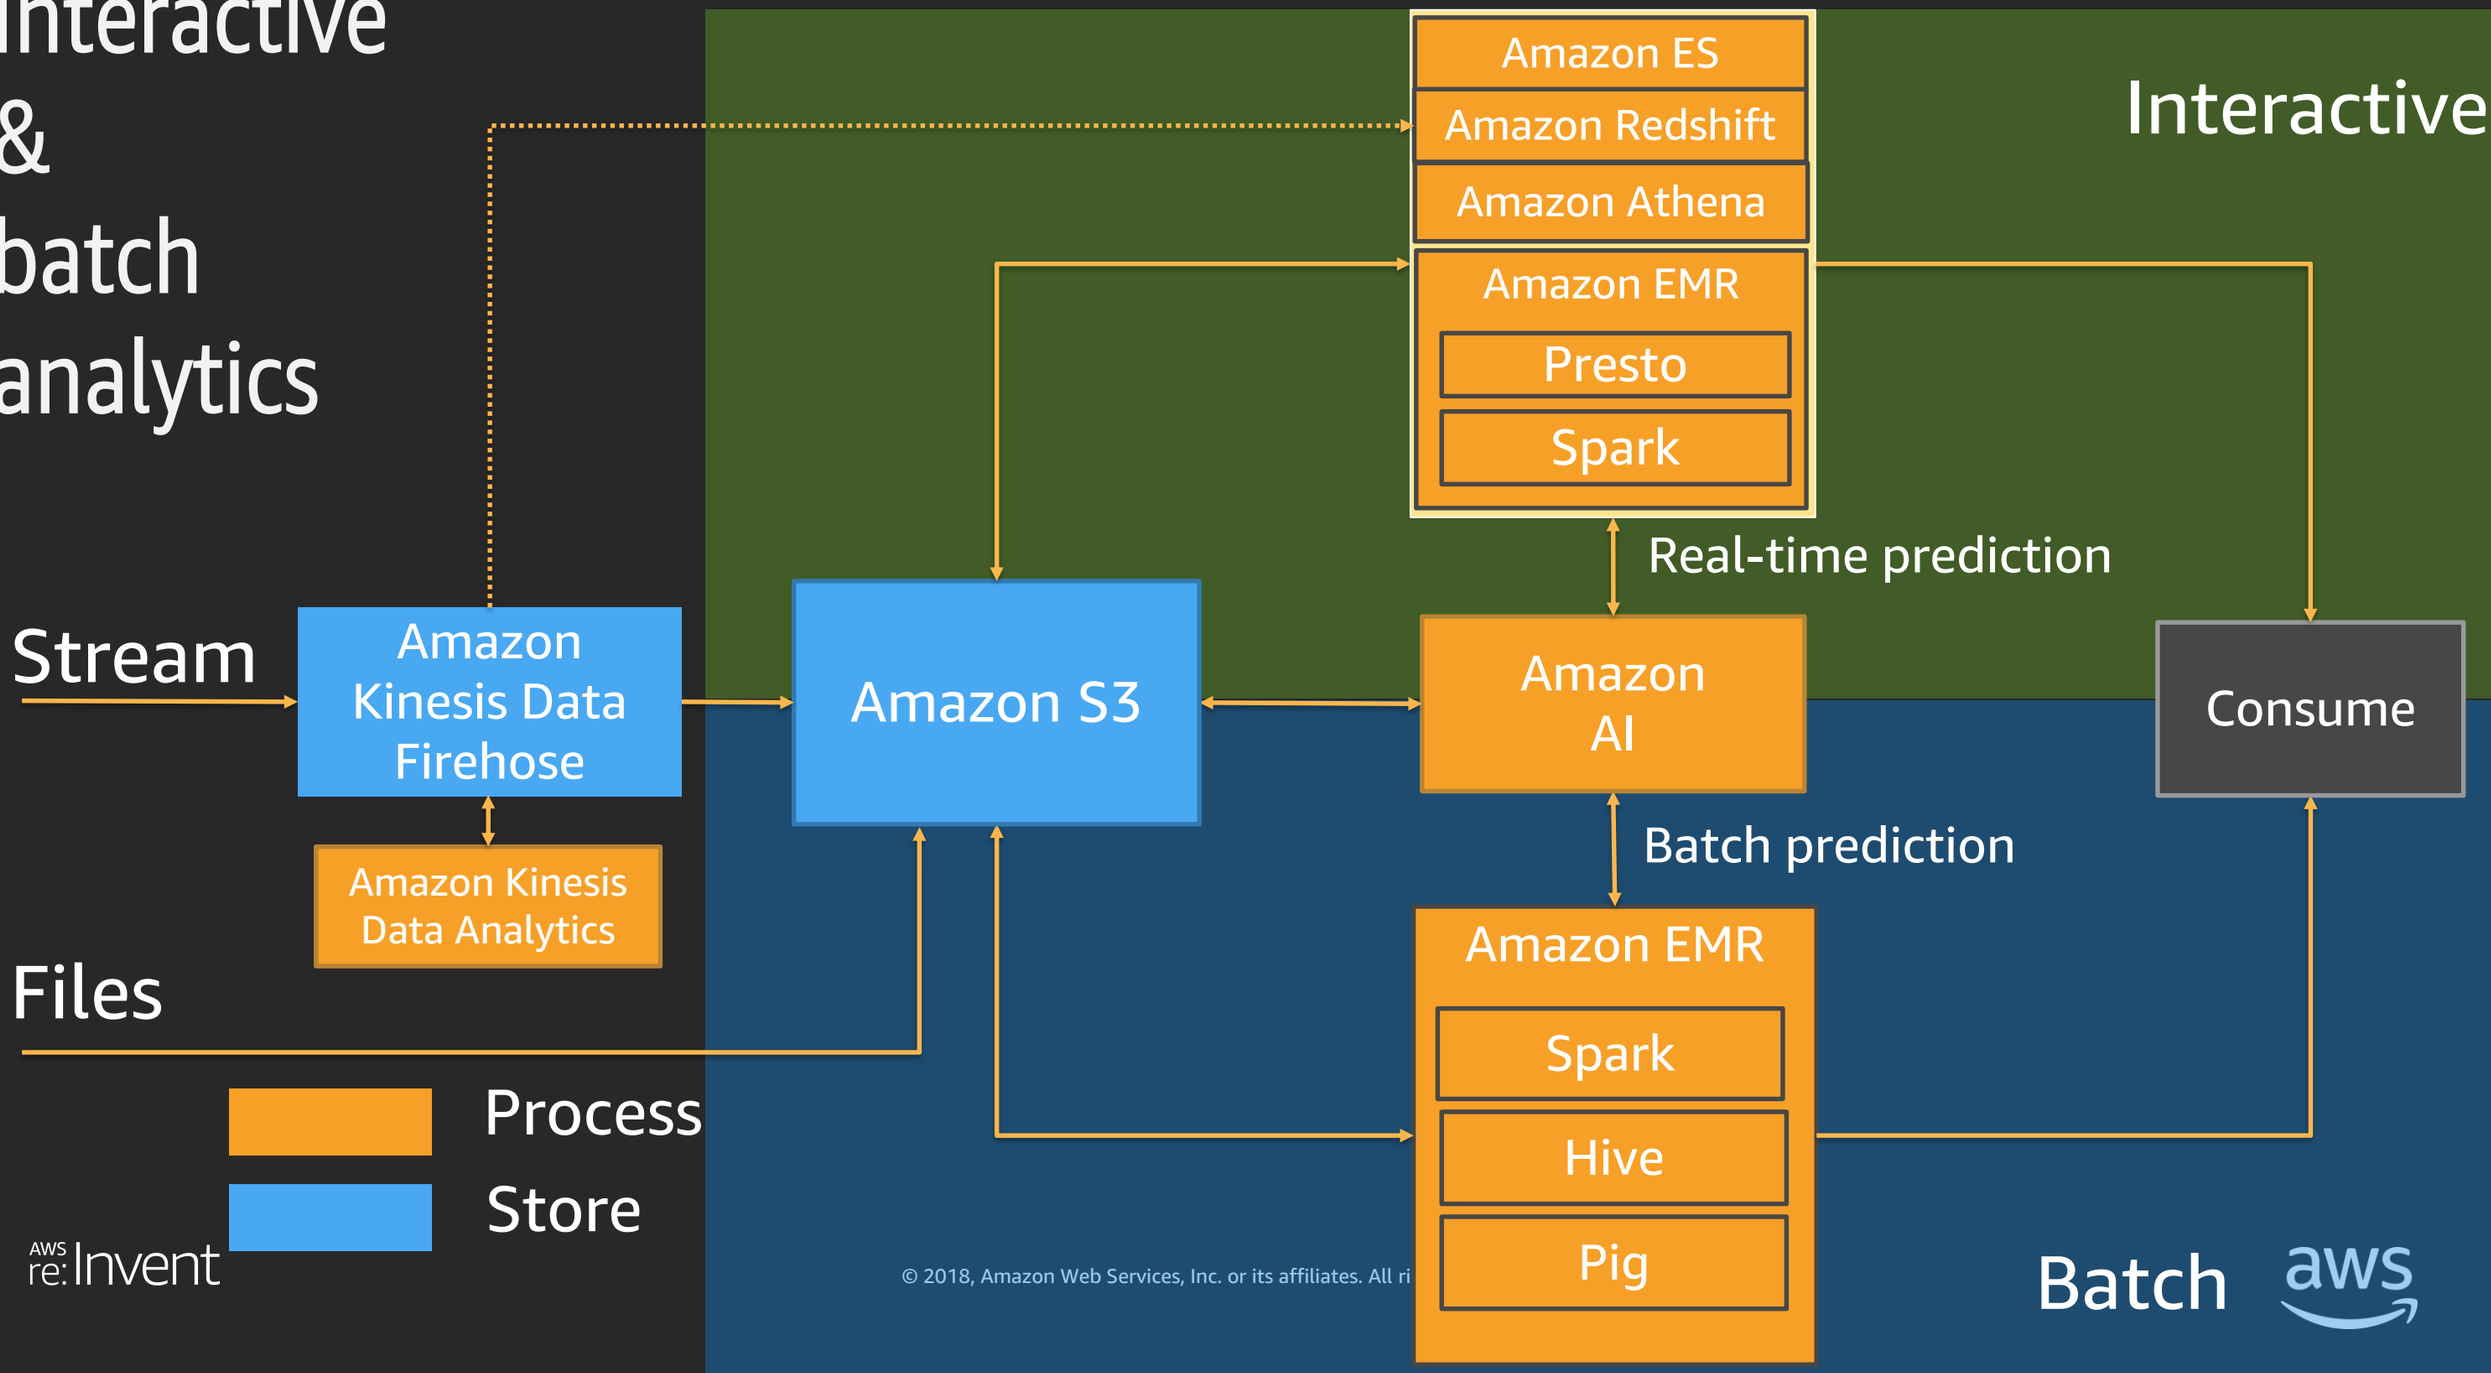

AWS re:Invent

aws

FINRA: Migrating to AWS

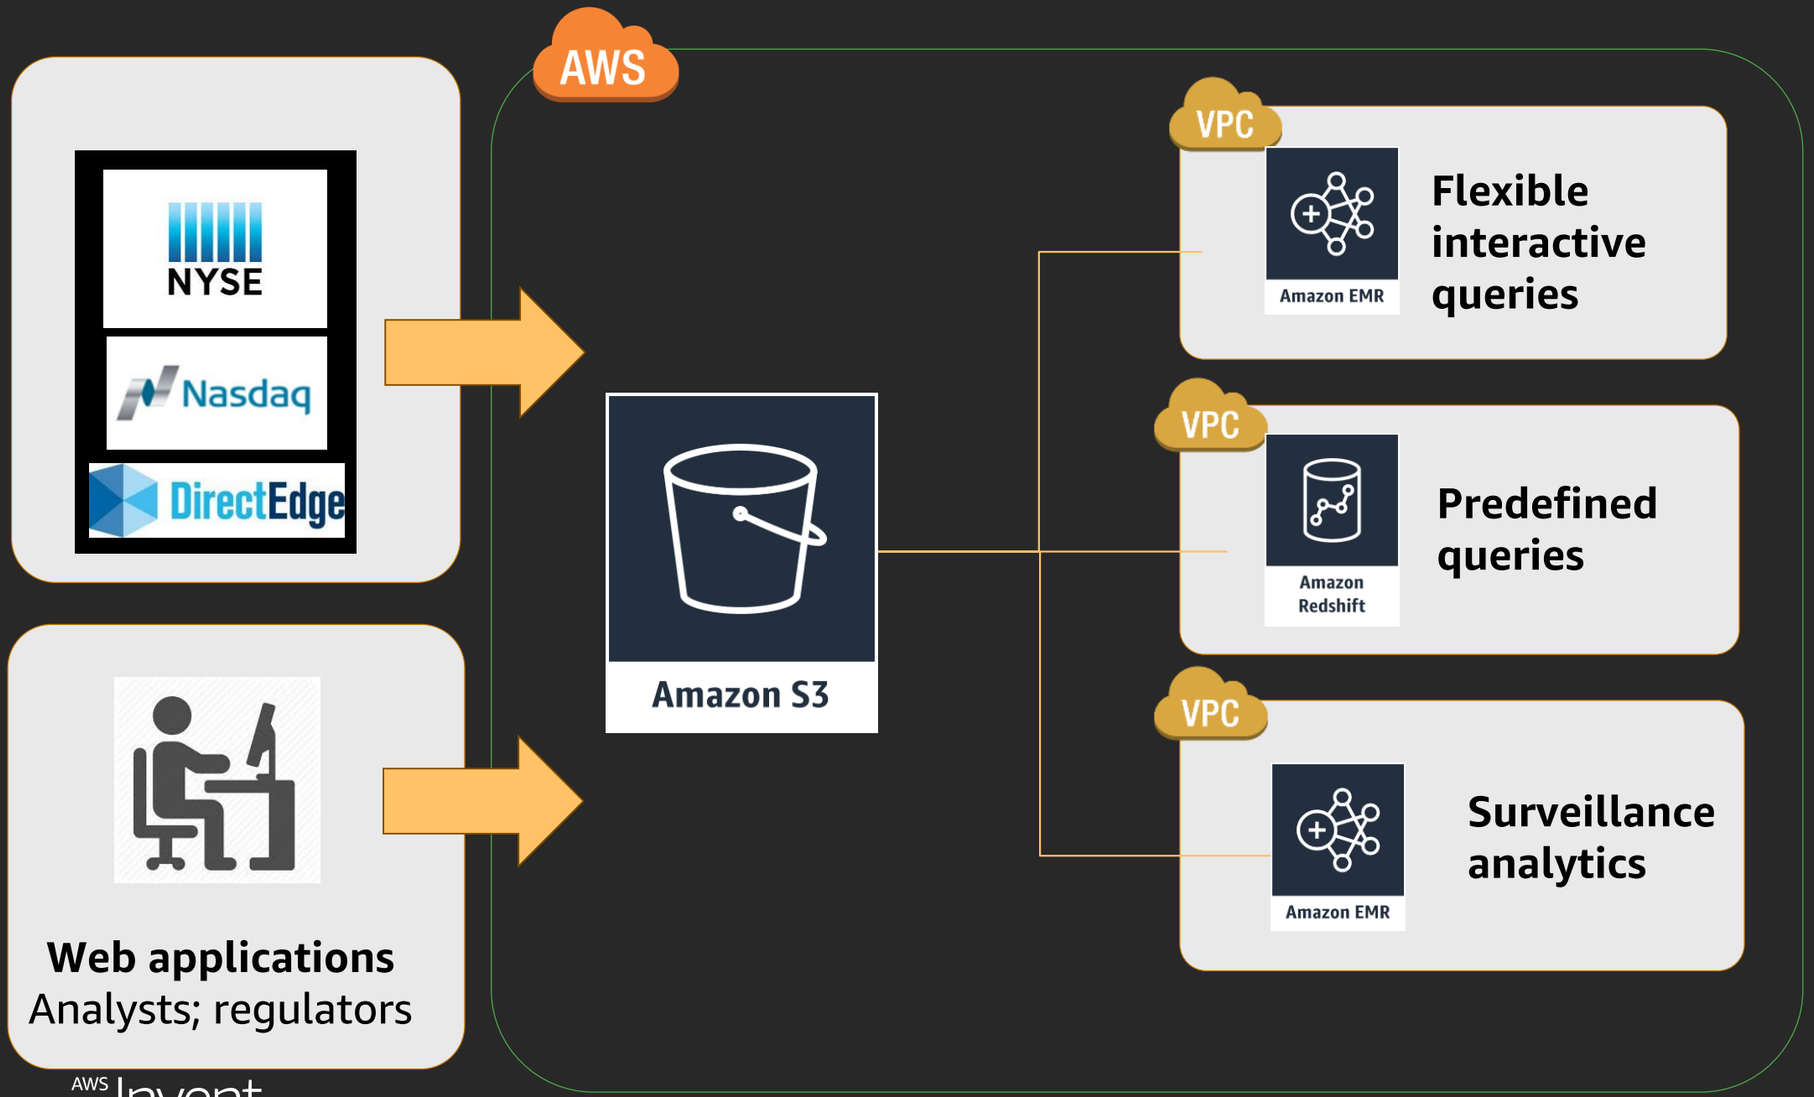

Petabytes of data generated on-premises, brought to AWS, and stored in S3

Thousands of analytical queries performed on Amazon EMR and Amazon Redshift

Stringent security requirements met by leveraging VPC, VPN, encryption at-rest and in-transit, AWS CloudTrail, and database auditing

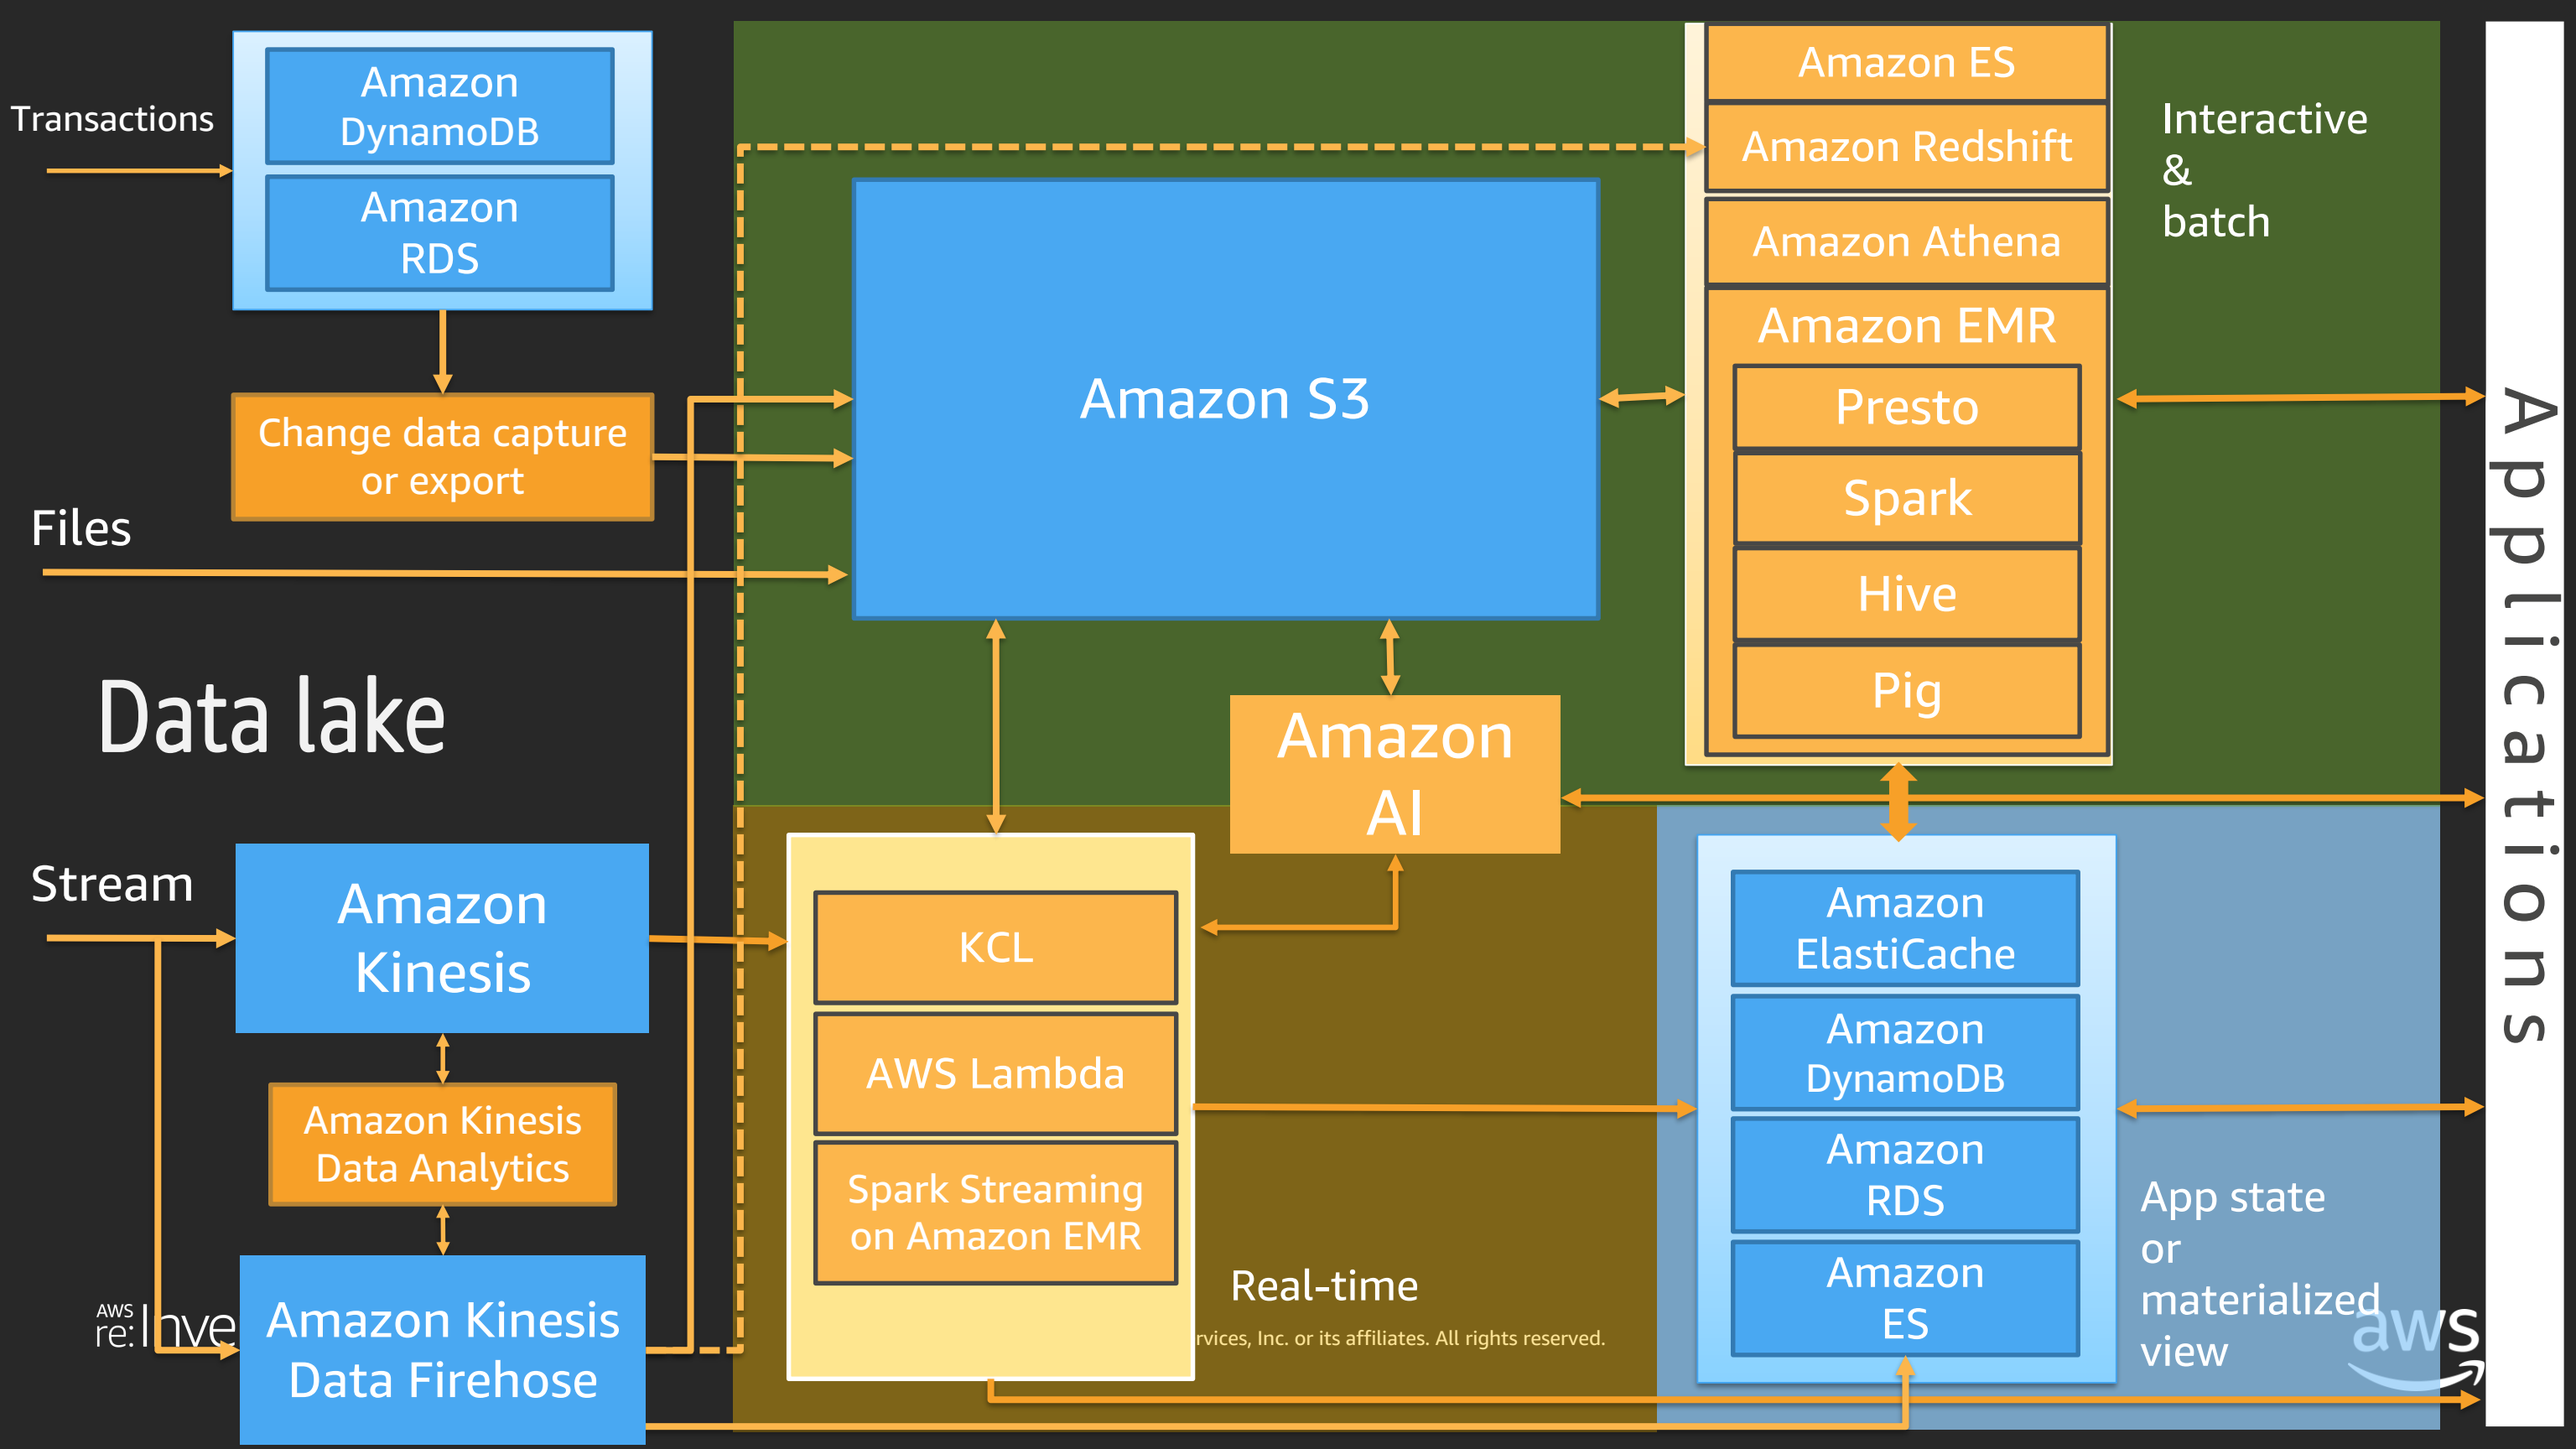

What about metadata?

- AWS Glue catalog
  - Hive Metastore compliant
  - Crawlers - Detect new data, schema, partitions
  - Search - Metadata discovery
  - Amazon Athena, Amazon EMR, and Amazon Redshift Spectrum compatible
- Hive Metastore (Presto, Spark, Hive, Pig)
  - Can be hosted on Amazon RDS

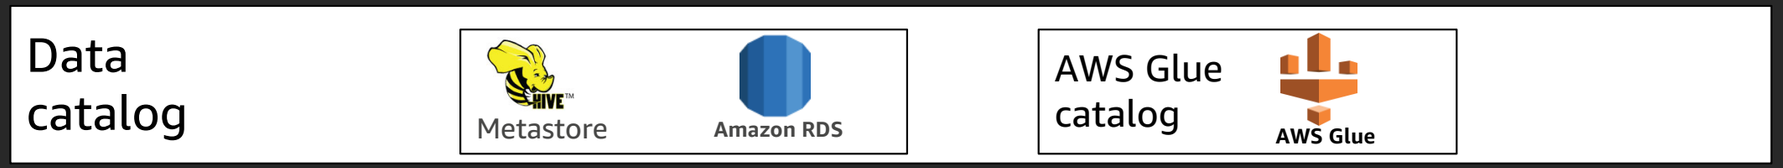

Demonstration—Data Lake Demonstration

Summary

Build decoupled systems

- Data → Store → Process → Store → Analyze → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Leverage AWS managed and serverless services

- Scalable/elastic, available, reliable, secure, no/low admin

Use log-centric design patterns

- Immutable logs, data lake, materialized views

Be cost-conscious

- Big data ≠ big cost

AI/ML enable your applications

Thank you!

Ben Snively

AWS re:Invent

aws

Please complete the session survey in the mobile app.

AWS re:Invent

aws